

Chapter 3: Processes

肖卿俊

办公室：九龙湖校区计算机楼212室

电邮：csqjxiao@seu.edu.cn

主页：<https://csqjxiao.github.io/PersonalPage>

电话：025-52091022



Chapter 3: Processes

- Process Concept
- Operations and APIs on Processes
- Process Scheduling
- Cooperating Processes
- Interprocess Communication
- Communication in Client-Server Systems



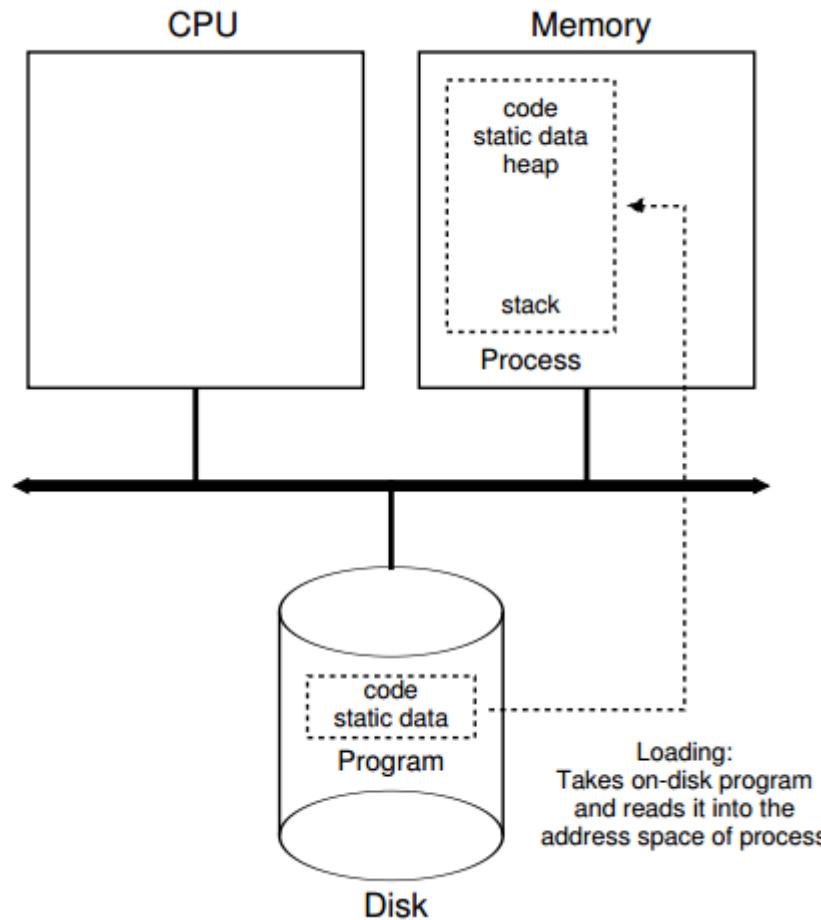
Process Concept

- An operating system executes a variety of programs:
 - ◆ Batch system – jobs
 - ◆ Time-shared systems – user programs or tasks
- Textbook uses the terms *job* and *process* almost interchangeably.
- Q: Why process, not program? What is a program?
- Process: running program
 - ◆ A program is lifeless, the OS makes it running (as a process).
 - ◆ A process can be viewed as a running program with machine states.

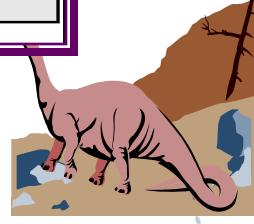
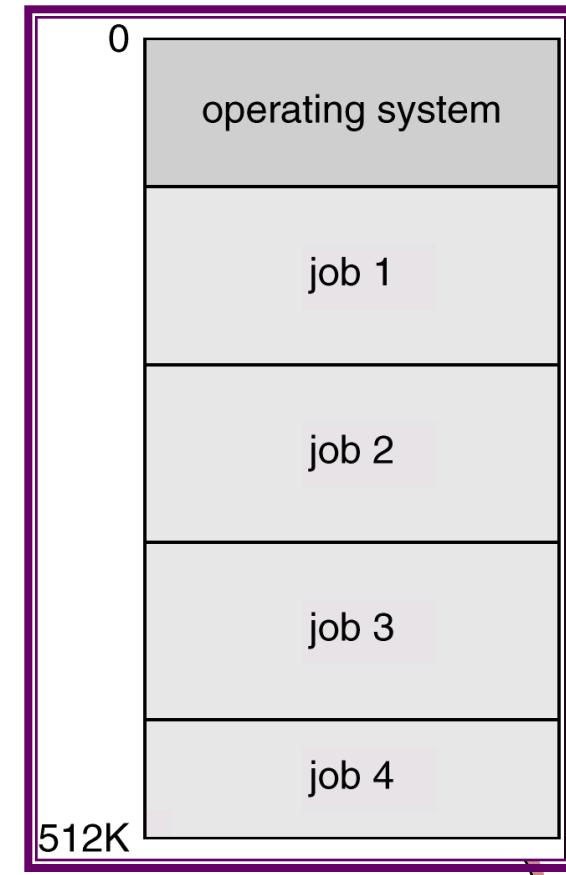




Loading into Memory: From Program To Process



Virtualizing the Memory





Processes

The Process Model

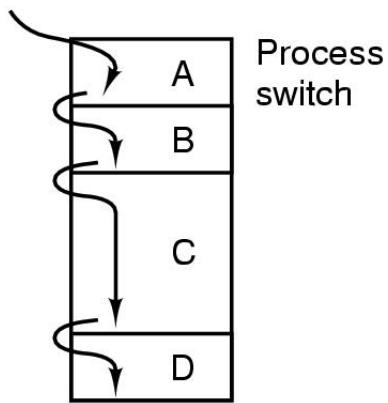
Virtualizing the CPU:

- By running one process, then stopping it and running another, and so forth.

An Example: Multiprogramming of four programs

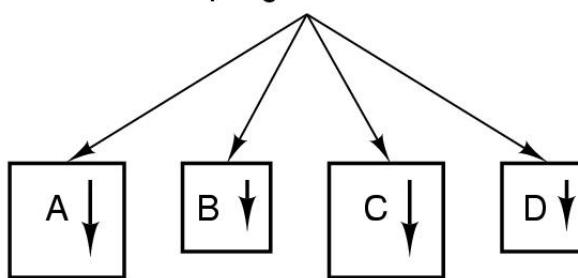
- Conceptual model of 4 independent, sequential processes
- Only one program active at any instant

One program counter

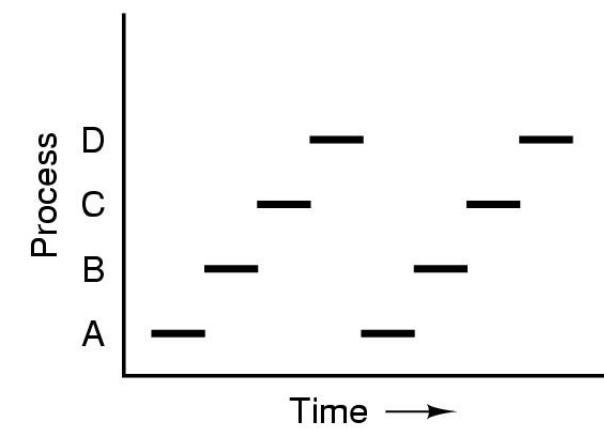


(a)

Four program counters



(b)



(c)





Process Concept (Cont.)

■ Process – a program in execution; process execution must progress in sequential fashion.

■ The running state of a process includes:

◆ Memory

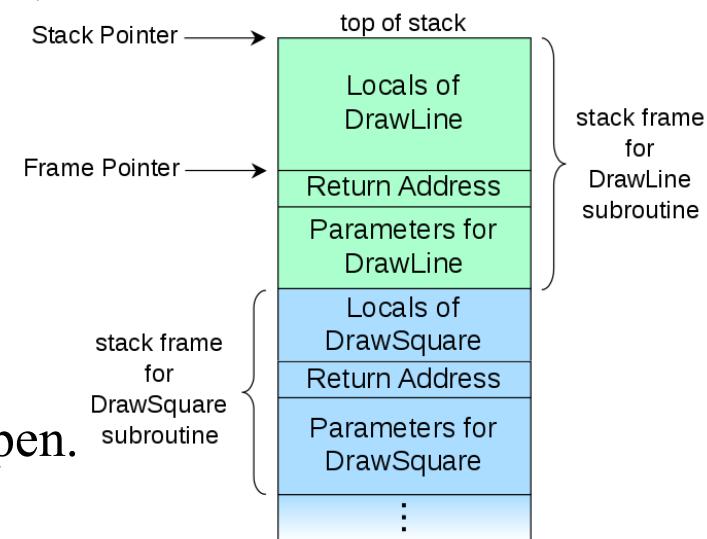
- ✓ **Address space**: Instructions and data.

◆ Registers

- ✓ **Program counter (PC) / instruction pointer (IP)**: current instruction.
- ✓ **Stack pointer, frame pointer**: management of stack for parameters, local variables, and return addresses.
- ✓ Contents of the processor's other registers

◆ I/O information

- ✓ A list of the files the process currently has open.





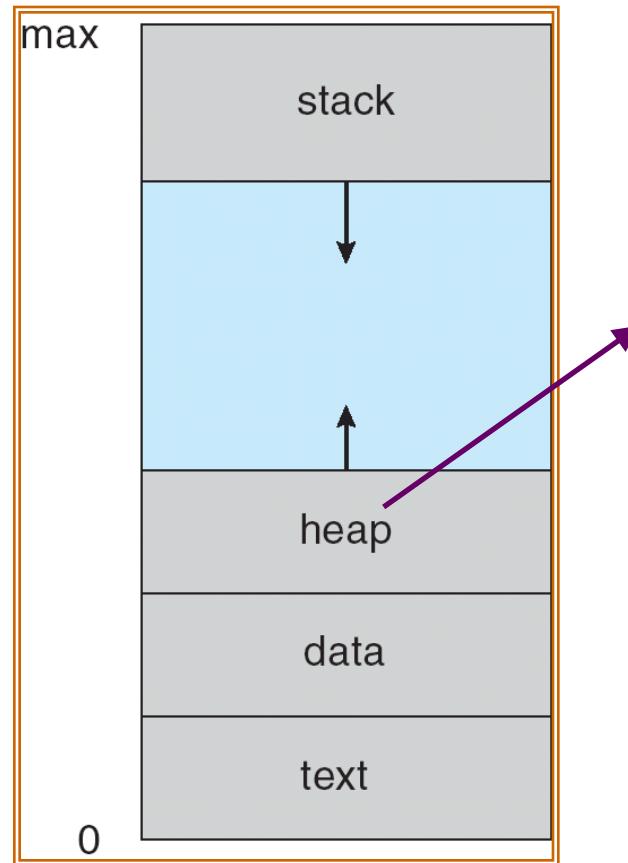
Process in Memory

```
//main.cpp
int a = 0; ← 数据段
char *p1; ← 数据段
main()
{
    int b; ← 栈
    char s[] = "abc"; ← 栈
    char *p2; ← 栈
    char *p3 = "123456"; ← 栈
    p1 = (char *)malloc(10); ← 堆
    p2 = (char *)malloc(20); ← 堆
}
```

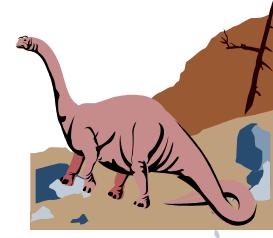




Process in Memory



Heap is memory that is dynamically allocated during process run time.





Process State

- As a process executes, it changes *state*
 - ◆ **new**: The process is being created.
 - ◆ **running**: Instructions are being executed.
 - ◆ **waiting**: The process is waiting for some event to occur.
 - ◆ **ready**: The process is waiting to be assigned to a processor.
 - ◆ **terminated**: The process has finished execution.

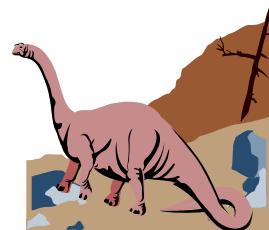
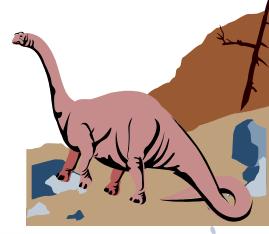
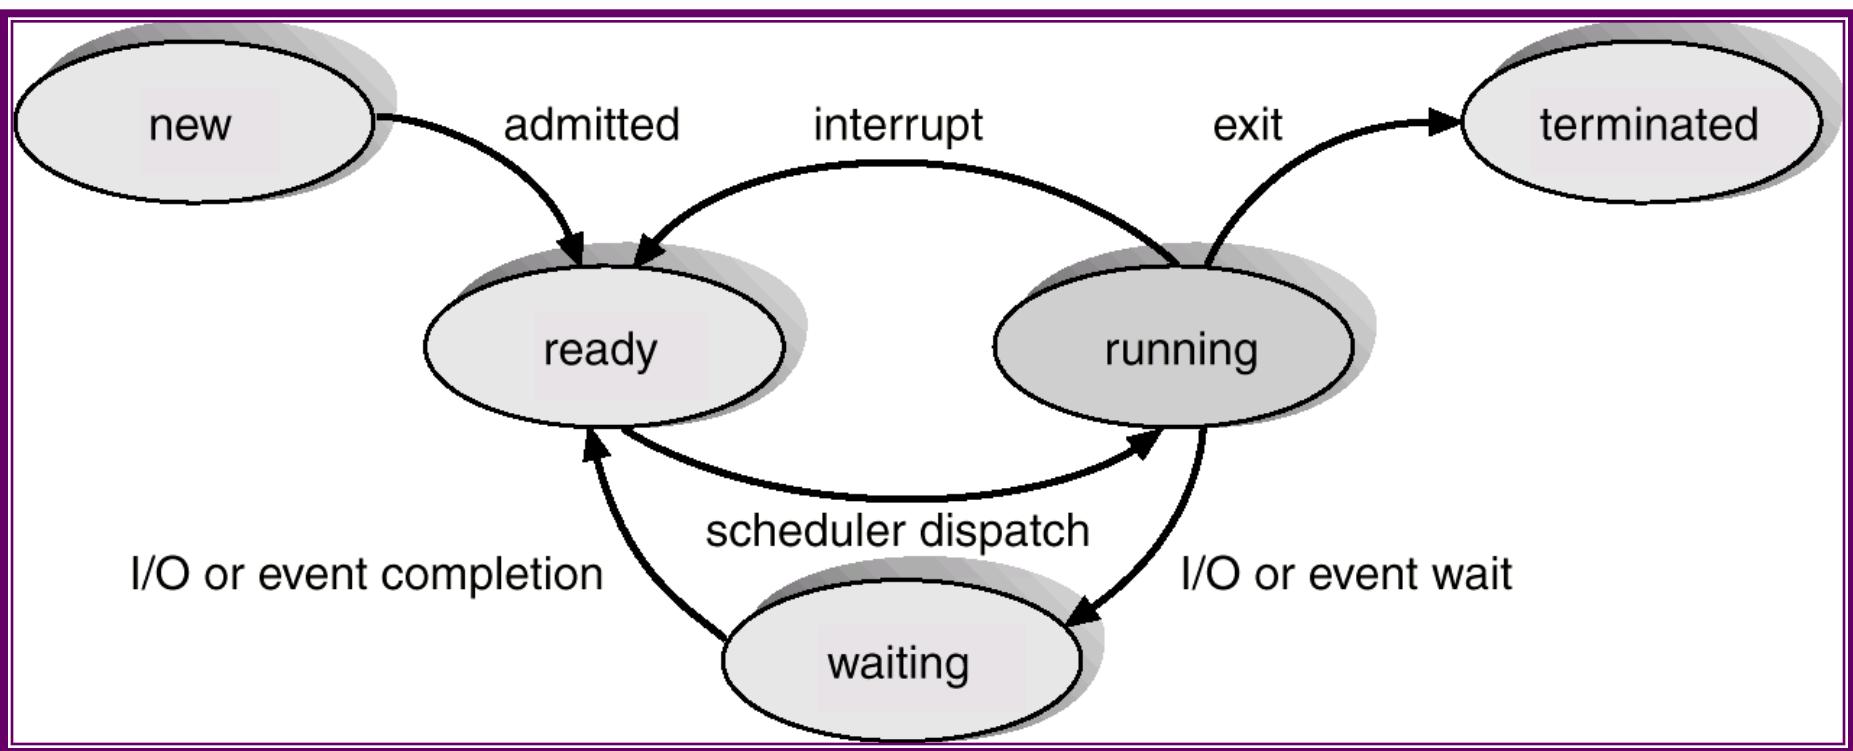




Diagram of Process State

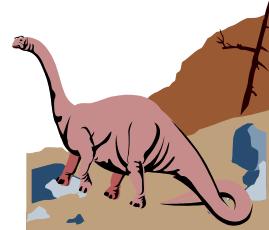




Tracing Process State

Time	Process ₀	Process ₁	Notes
1	Running	Ready	
2	Running	Ready	
3	Running	Ready	Process ₀ initiates I/O
4	Blocked	Running	Process ₀ is blocked, so Process ₁ runs
5	Blocked	Running	
6	Blocked	Running	
7	Ready	Running	I/O done
8	Ready	Running	Process ₁ now done
9	Running	—	
10	Running	—	Process ₀ now done

Figure 4.4: Tracing Process State: CPU and I/O





Discussion

- Q1: Draw on the blackboard the Diagram of Process State

- Q2: 下列哪一种情况不会引起进程之间的切换?
 - ◆ A. 进程调用本程序中定义的函数进行计算
 - ◆ B. 进程处理I/O请求
 - ◆ C. 进程创建子进程并等待子进程结束
 - ◆ D. 产生中断





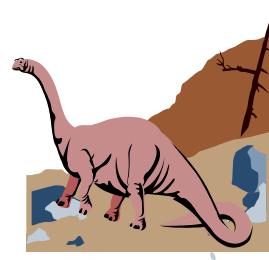
Data Structure

■ OS is a program, so it has some key data structures that track the state of each process.

- ◆ Process lists for all ready / running / waiting processes

■ An example: xv6 kernel

- ◆ types of information an OS needs to track processes





```
// the registers xv6 will save and restore
// to stop and subsequently restart a process
struct context {
    int eip;
    int esp;
    int ebx;
    int ecx;
    int edx;
    int esi;
    int edi;
    int ebp;
};

// the different states a process can be in
enum proc_state { UNUSED, EMBRYO, SLEEPING,
                  RUNNABLE, RUNNING, ZOMBIE };

// the information xv6 tracks about each process
// including its register context and state
struct proc {
    char *mem;                      // Start of process memory
    uint sz;                        // Size of process memory
    char *kstack;                   // Bottom of kernel stack
                                    // for this process
    enum proc_state state;          // Process state
    int pid;                        // Process ID
    struct proc *parent;            // Parent process
    void *chan;                     // If non-zero, sleeping on chan
    int killed;                     // If non-zero, have been killed
    struct file *ofile[NFILE];     // Open files
    struct inode *cwd;              // Current directory
    struct context context;         // Switch here to run process
    struct trapframe *tf;           // Trap frame for the
                                    // current interrupt
};

Operating System( );
```

All registers

Memory:
address space

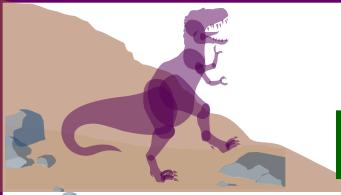
Stack

Process state

Process ID

I/O information

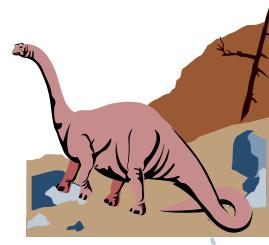




Process Control Block (PCB)

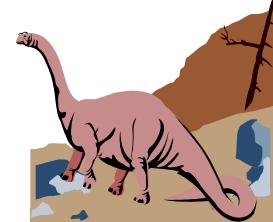
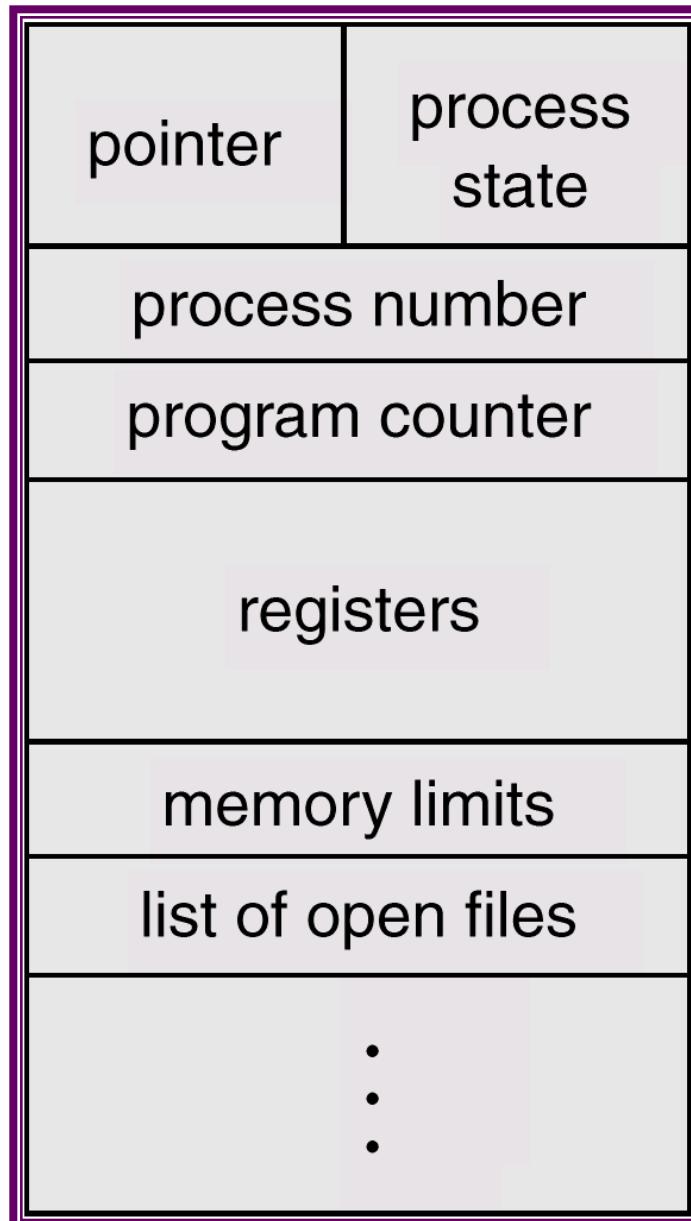
■ Information associated with each process.

- ◆ Process state
- ◆ Program counter
- ◆ CPU registers
- ◆ CPU scheduling information
- ◆ Memory-management information
- ◆ Accounting information
- ◆ File usage and I/O status information





Process Control Block (PCB)

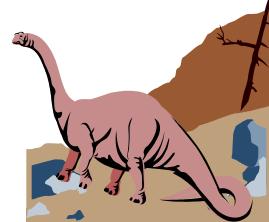




Context Switch

■ What is a process context?

- ◆ The *context* of a process includes the values of CPU registers, the process state, the program counter, and other memory/file management information.

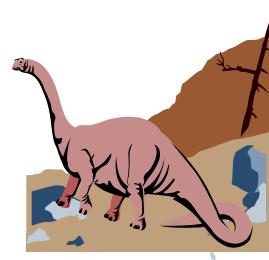




Context Switch (cont.)

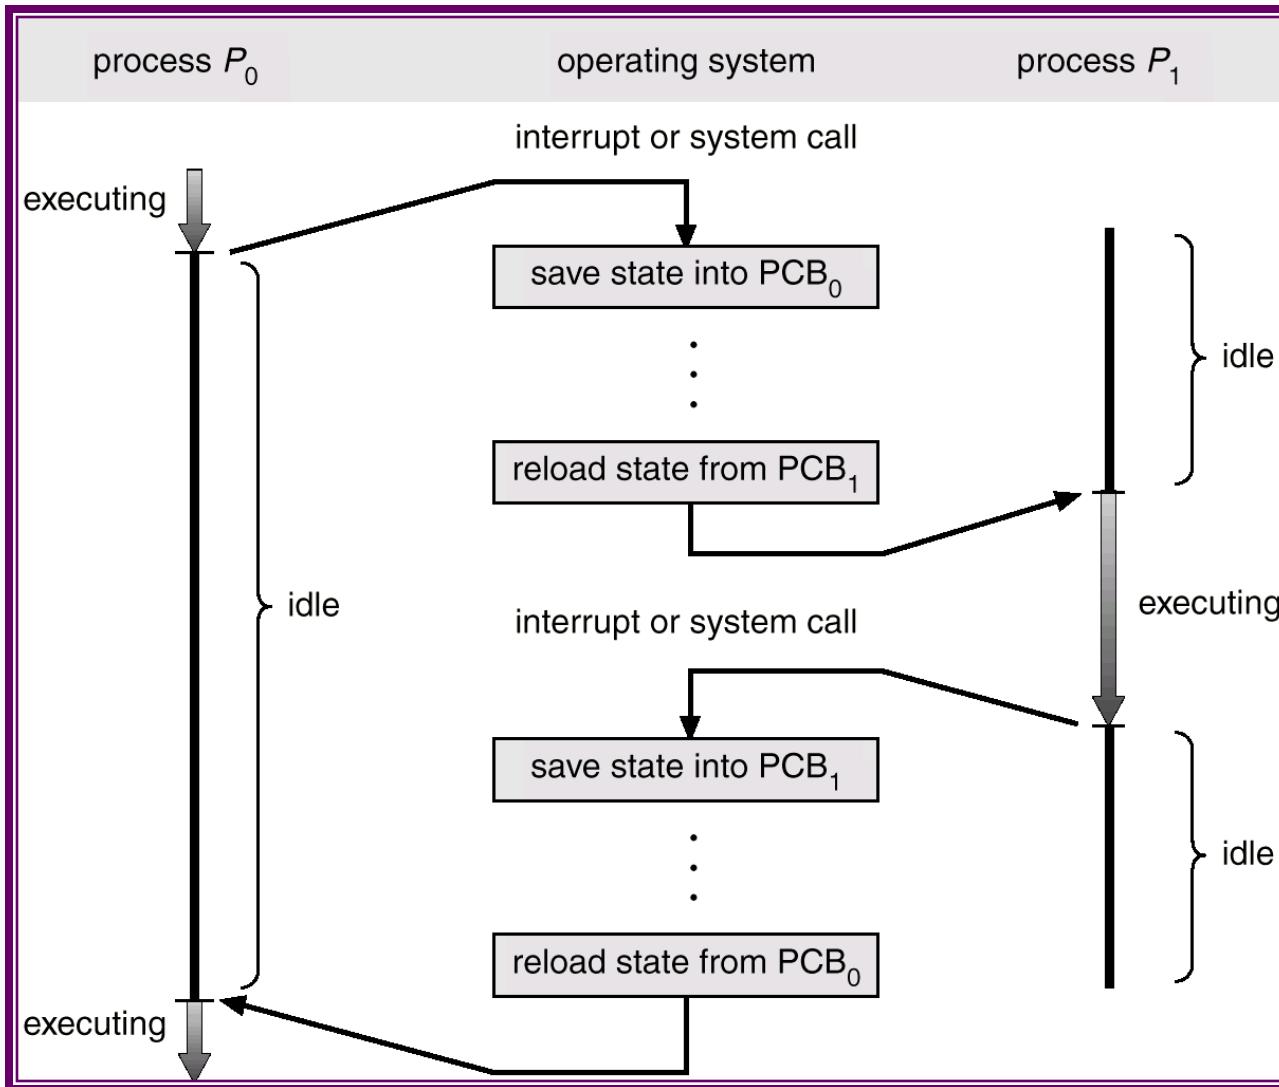
■ What is a context switch?

- ◆ After the CPU scheduler selects a process (from the *ready queue*) and before allocates CPU to it, the CPU scheduler must
 - ✓ save the ***context*** of the currently running process,
 - ✓ put it into a queue,
 - ✓ load the ***context*** of the selected process, and
 - ✓ let it run.





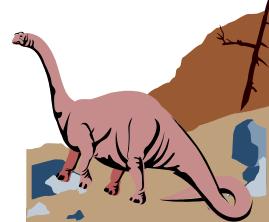
CPU Switch From Process to Process





Context Switch (Cont.)

- When CPU switches to another process, the system must save the state of the old process and load the saved state for the new process.
- Context-switch time is overhead; the system does no useful work while switching.
- Time dependent on hardware support.





Chapter 3: Processes

- Process Concept
- Operations and APIs on Processes
- Process Scheduling
- Cooperating Processes
- Interprocess Communication
- Communication in Client-Server Systems



Process Creation

■ Parent process create children processes, which, in turn create other processes, forming a tree of processes.

■ Resource sharing

- ◆ Parent and children share all resources.
- ◆ Children share subset of parent's resources.
- ◆ Parent and child share no resources.

■ Execution

- ◆ Parent and children execute concurrently.
- ◆ Parent waits until children terminate.





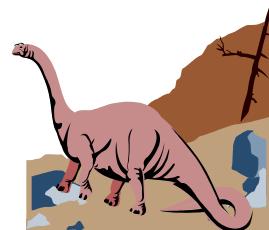
Process Creation (Cont.)

■ Address space

- ◆ Child duplicate of parent.
- ◆ Child has a program loaded into it.

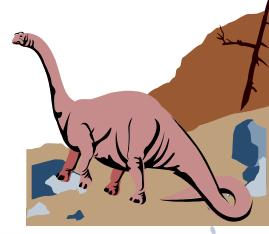
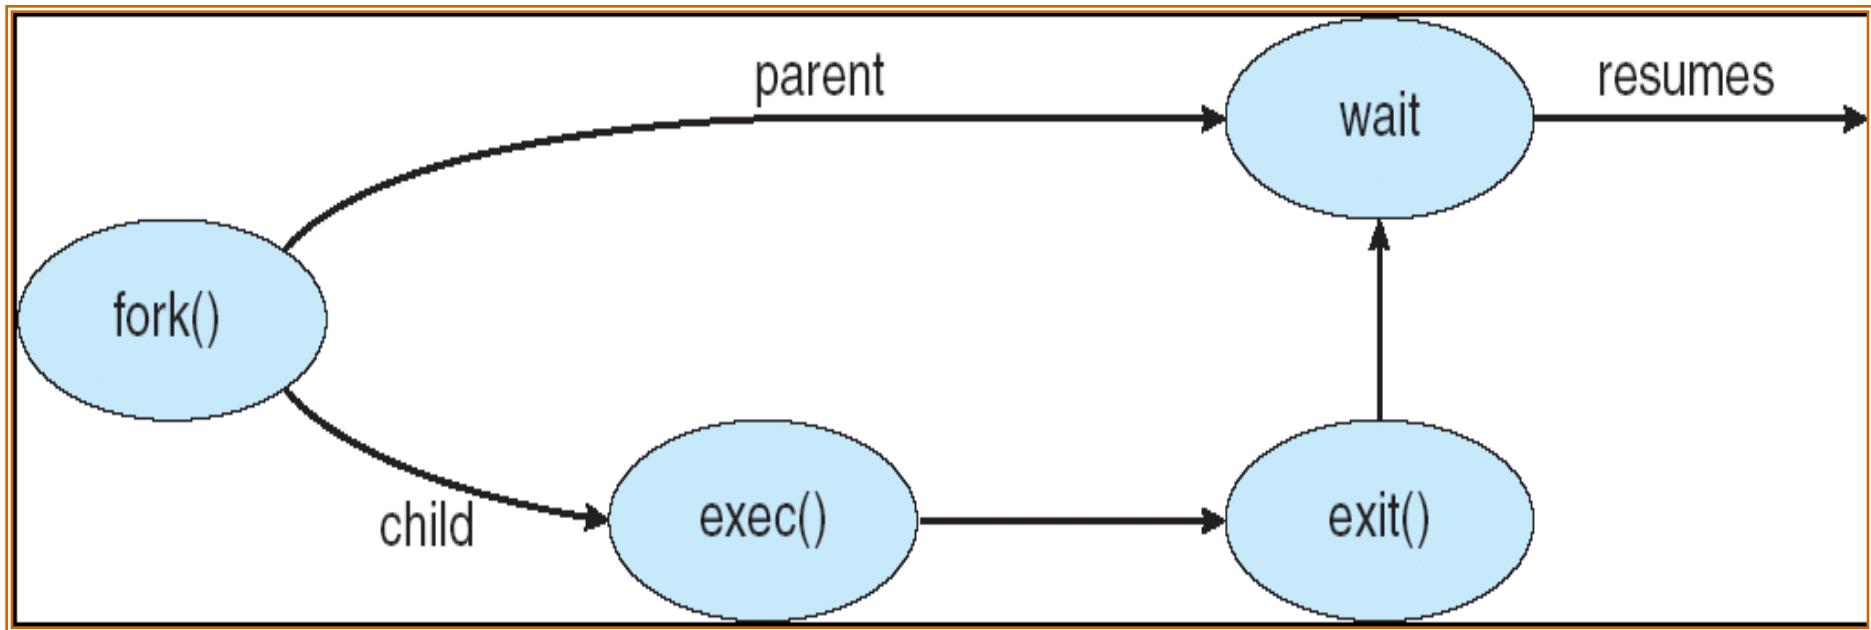
■ UNIX examples

- ◆ **fork** system call creates new process
- ◆ **exec** system call used after a **fork** to replace the process' memory space with a new program.





Process Creation (UNIX)





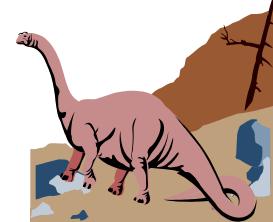
The fork() System Call

```
1 #include <stdio.h>
2 #include <stdlib.h>
3 #include <unistd.h>
4
5 int
6 main(int argc, char *argv[])
7 {
8     printf("hello world (pid:%d)\n", (int) getpid());
9     int rc = fork();
10    if (rc < 0) {           // fork failed; exit
11        fprintf(stderr, "fork failed\n");
12        exit(1);
13    } else if (rc == 0) {   // child (new process)
14        printf("hello, I am child (pid:%d)\n", (int) getpid());
15    } else {                // parent goes down this path (main)
16        printf("hello, I am parent of %d (pid:%d)\n",
17               rc, (int) getpid());
18    }
19    return 0;
20 }
```

— - - - -

```
prompt> ./p1
hello world (pid:29146)
hello, I am parent of 29147 (pid:29146)
hello, I am child (pid:29147)
prompt>
```

ODD?





The fork() System Call

- The process that is created by using the fork() system call is an (almost) exact **copy** of the calling process.

- Discussion:

what is the output?

```
int rc = fork();  
if (rc < 0) {  
    printf("A");  
    exit(1);  
} else if (rc == 0) {  
    printf("B");  
} else {  
    printf("C");  
}  
return 0;
```





The fork() System Call

- Discussion: What is the output if we add a loop command before the screen print command?

```
int
main(int argc, char *argv[])
{
    printf("hello world (pid:%d)\n", (int) getpid());
    int rc = fork();
    if (rc < 0) {
        // fork failed; exit
        fprintf(stderr, "fork failed\n");
        exit(1);
    } else if (rc == 0) {
        // child (new process)
        int sum = 0;
        for (int i = 0; i < 1000000000; i++)
            sum += i;
        printf("hello, I am child (pid:%d)\n", (int) getpid());
    } else {
        // parent goes down this path (original process)
        int sum = 0;
        for (int i = 0; i < 1000000000; i++)
            sum += i;
        printf("hello, I am parent of %d (pid:%d)\n",
               rc, (int) getpid());
    }
    return 0;
}
```

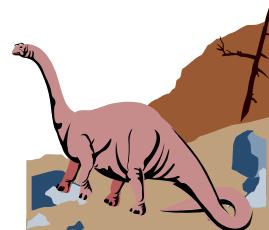




The fork() System Call

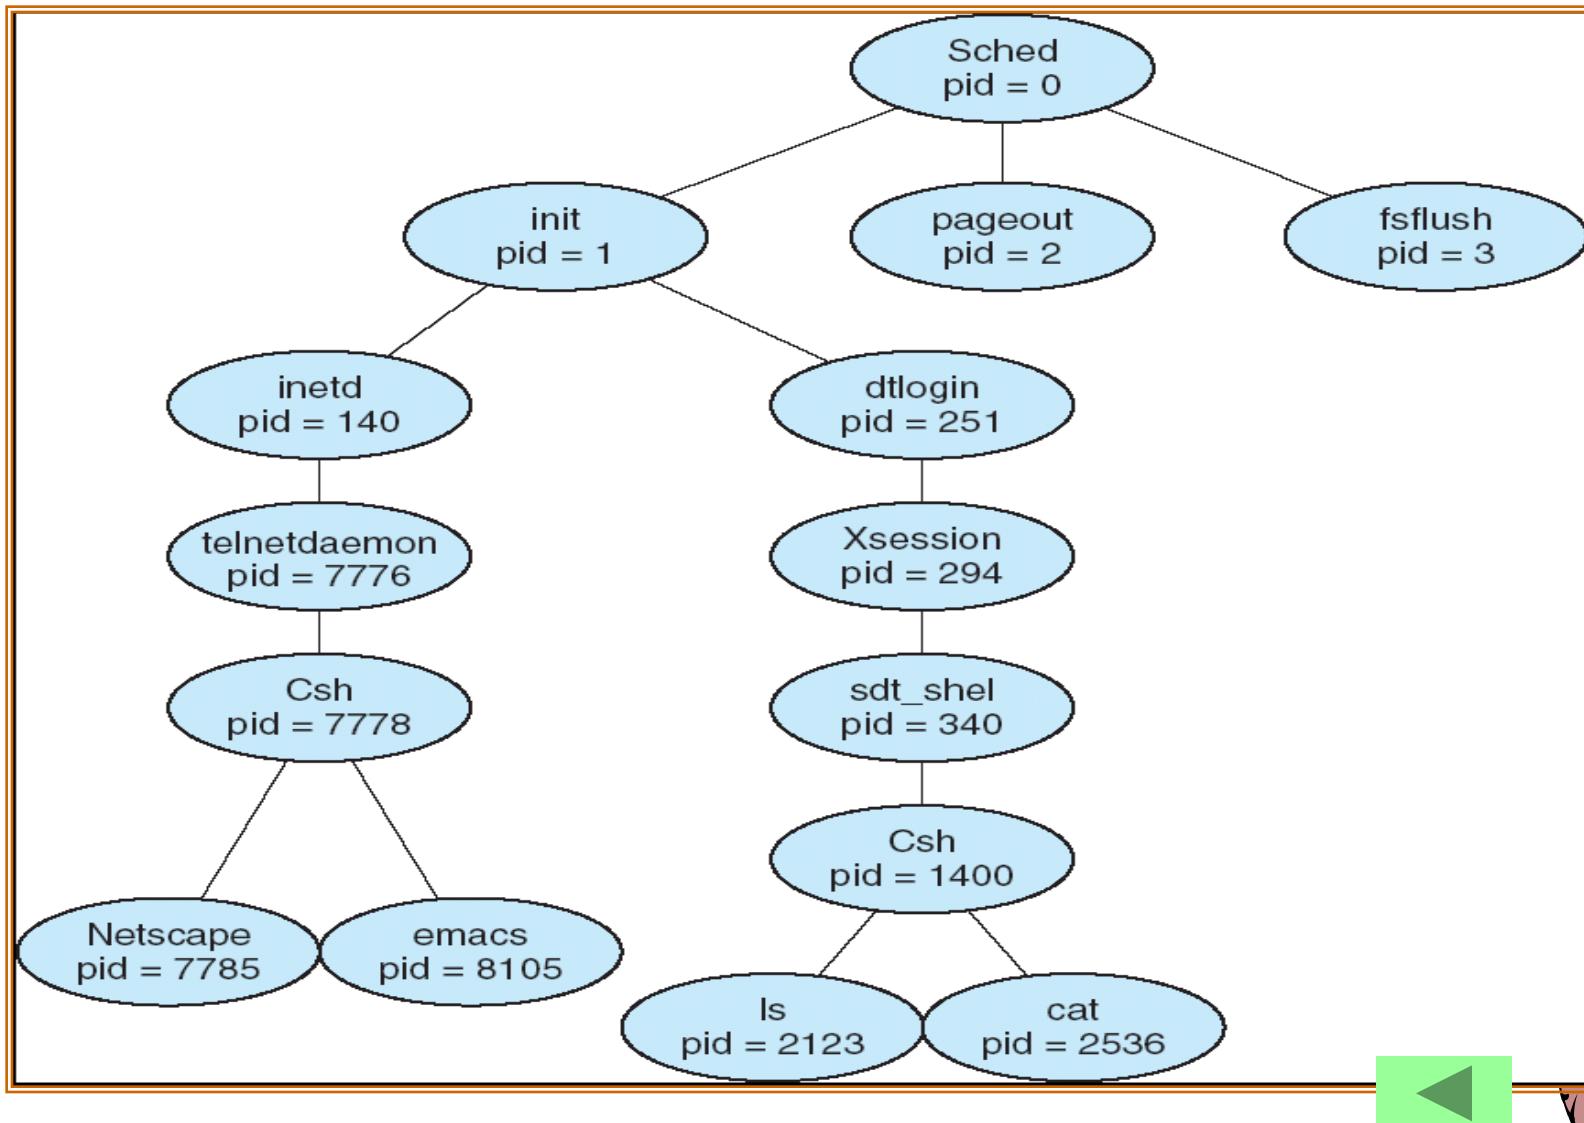
```
Qingjuns-MacBook-Pro:OSC3_code_cpu-api csqjxiao$ ./p1-2
hello world (pid:43349)
hello, I am parent of 43350 (pid:43349)
hello, I am child (pid:43350)
Qingjuns-MacBook-Pro:OSC3_code_cpu-api csqjxiao$ ./p1-2
hello world (pid:43352)
hello, I am child (pid:43353)
hello, I am parent of 43353 (pid:43352)
Qingjuns-MacBook-Pro:OSC3_code_cpu-api csqjxiao$ █
```

■ Discussion: why not deterministic?





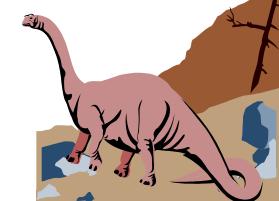
Processes Tree on Solaris





Process Termination

- Process executes last statement and asks the operating system to delete it (**exit**).
 - ◆ Output data from child to parent (via **wait**).
 - ◆ Process' resources are deallocated by OS.
- Parent may terminate execution of children processes (**abort**).
 - ◆ Child has exceeded allocated resources.
 - ◆ Task assigned to child is no longer required.
 - ◆ Parent is exiting.
 - ✓ Operating system does not allow child to continue if its parent terminates.
 - ✓ Cascading termination.





The wait() System Call

```
1 #include <stdio.h>
2 #include <stdlib.h>
3 #include <unistd.h>
4 #include <sys/wait.h>
5
6 int
7 main(int argc, char *argv[])
8 {
9     printf("hello world (pid:%d)\n", (int) getpid());
10    int rc = fork();
11    if (rc < 0) {           // fork failed; exit
12        fprintf(stderr, "fork failed\n");
13        exit(1);
14    } else if (rc == 0) {   // child (new process)
15        printf("hello, I am child (pid:%d)\n", (int) getpid());
16    } else {               // parent goes down this path (main)
17        int wc = wait(NULL);
18        printf("hello, I am parent of %d (wc:%d) (pid:%d)\n",
19                rc, wc, (int) getpid());
20    }
21    return 0;
22 }
```

parent waits for child process to finish

```
prompt> ./p2
hello world (pid:29266)
hello, I am child (pid:29267)
hello, I am parent of 29267 (wc:29267) (pid:29266)
prompt>
```

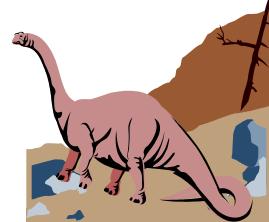




The exec() System Call

- The process that is created by using the exec() system call can be a different program.

- Some details in exec()
 - ◆ It does not create a new process; rather, it transforms the currently running program into a different running program.

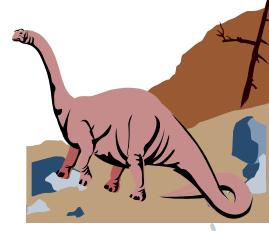




The exec() System Call

```
1 #include <stdio.h>
2 #include <stdlib.h>
3 #include <unistd.h>
4 #include <string.h>
5 #include <sys/wait.h>
6
7 int
8 main(int argc, char *argv[])
9 {
10     printf("hello world (pid:%d)\n", (int) getpid());
11     int rc = fork();
12     if (rc < 0) {           // fork failed; exit
13         fprintf(stderr, "fork failed\n");
14         exit(1);
15     } else if (rc == 0) {   // child (new process)
16         printf("hello, I am child (pid:%d)\n", (int) getpid());
17         char *myargs[3];
18         myargs[0] = strdup("wc");    // program: "wc" (word count)
19         myargs[1] = strdup("p3.c"); // argument: file to count
20         myargs[2] = NULL;          // marks end of array
21         execvp(myargs[0], myargs); // runs word count
22         printf("this shouldn't print out");
23     } else {               // parent goes down this path (main)
24         int wc = wait(NULL);
25         printf("hello, I am parent of %d (wc:%d) (pid:%d)\n",
26                 rc, wc, (int) getpid());
27     }
28     return 0;
29 }
```

prompt> ./p3
hello world (pid:29383)
hello, I am child (pid:29384)
 29 107 1030 p3.c
hello, I am parent of 29384 (wc:29384) (pid:29383)





Review

■ Process creation APIs

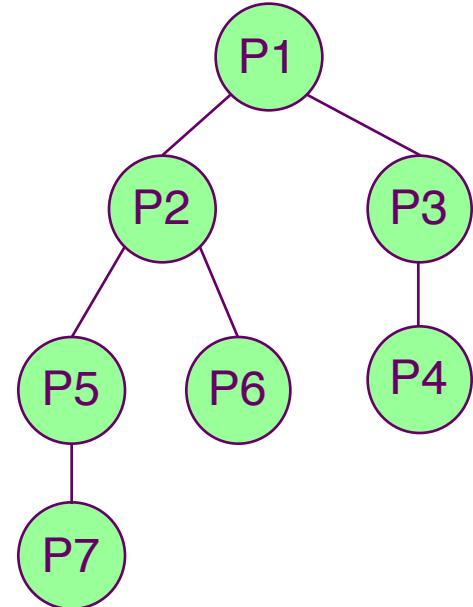
- ◆ fork()
 - ◆ wait()
 - ◆ exec()
-
- ◆ What are the differences?



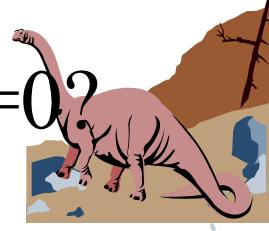
Quiz

■ Consider the following C program. Guess how many lines of output will be printed.

```
int main(int argc, char * argv[])
{
    int i, id1, id2;
    for (i = 1; i < 2; i++) {
        id1 = fork();
        id2 = fork();
        if (id1 == 0 || id2 == 0) fork();
    }
    printf("I am %d\n", getpid());
}
```



■ What if we change the initial value $i=1$ to $i=0$?





Chapter 3: Processes

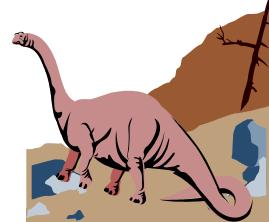
- Process Concept
- Operations and APIs on Processes
- **Process Scheduling**
- Cooperating Processes
- Interprocess Communication
- Communication in Client-Server Systems



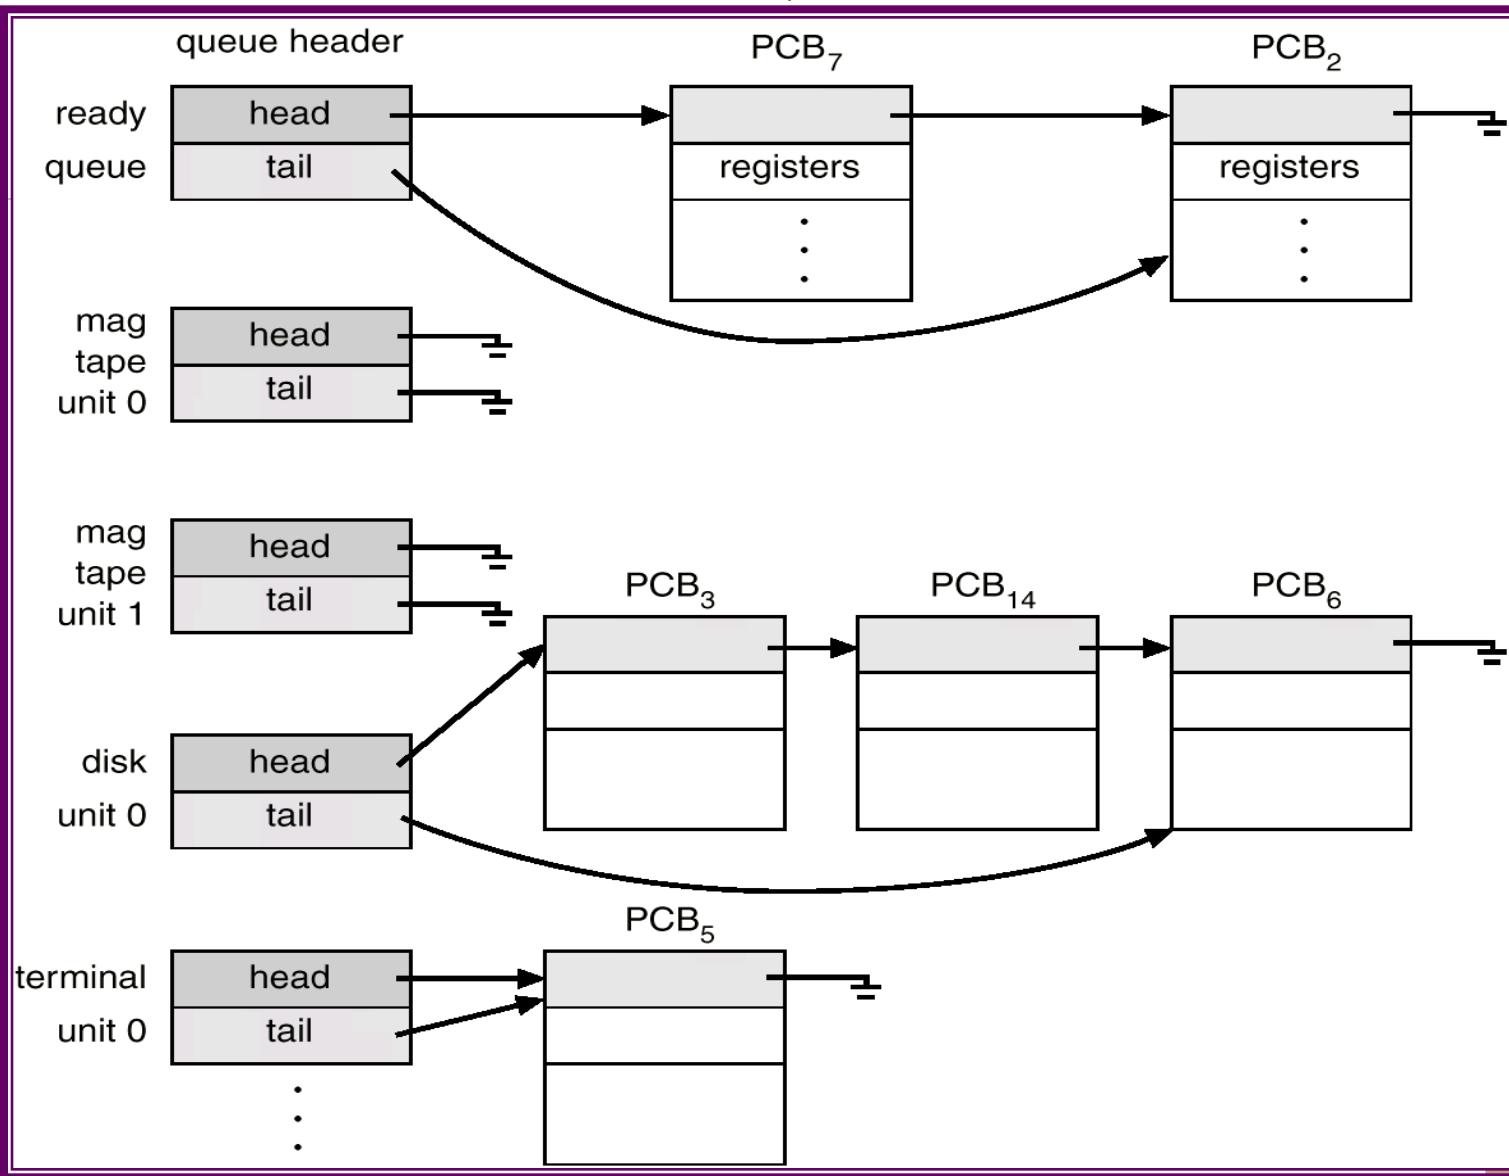


Process Scheduling Queues

- **Job queue** – set of all processes in the system.
- **Ready queue** – set of all processes residing in main memory, ready and waiting to execute.
- **Device queues** – set of processes waiting for an I/O device.
- Process migration between the various queues.

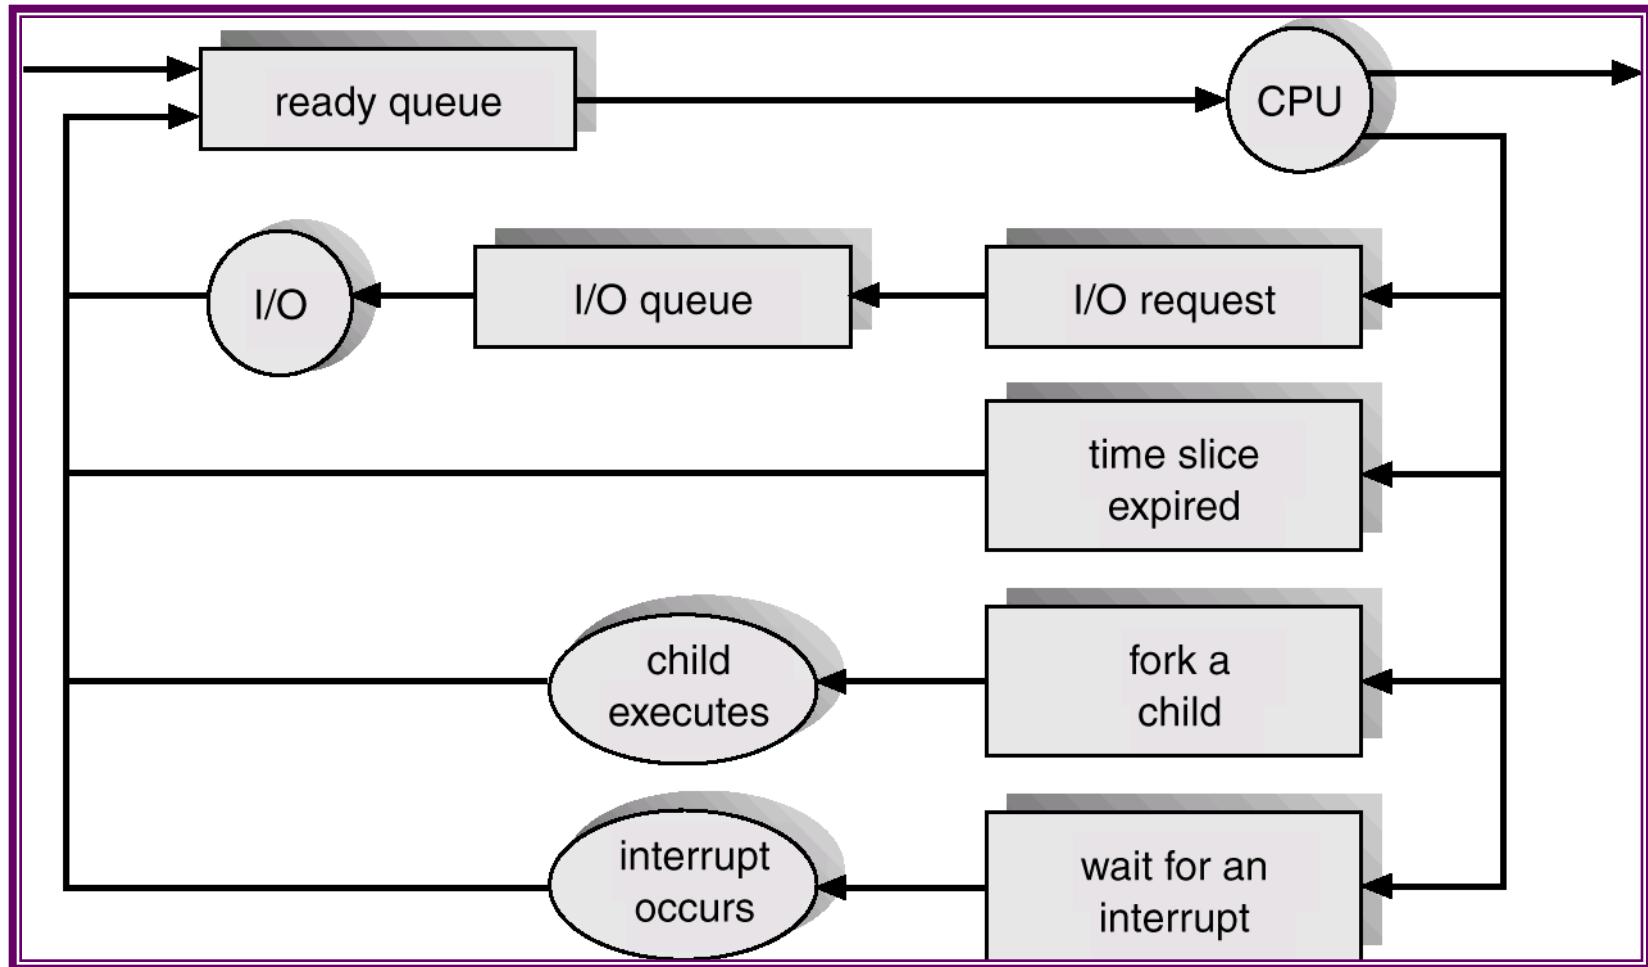


Ready Queue And Various I/O Device Queues





Representation of Process Scheduling





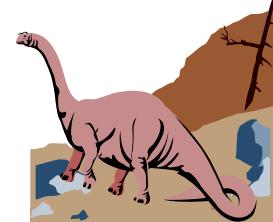
Schedulers

- Long-term scheduler (or job scheduler) – selects which processes should be loaded into memory for execution.
- Short-term scheduler (or CPU scheduler) – selects which process should be executed next and allocates CPU.



Schedulers (Cont.)

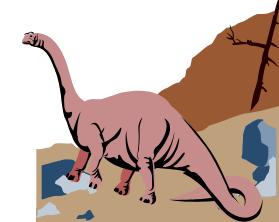
- Short-term scheduler is invoked very frequently (milliseconds) \Rightarrow (must be fast).
- Long-term scheduler is invoked very infrequently (seconds, minutes) \Rightarrow (may be slow).
- The _____ scheduler controls the *degree of multiprogramming*.
 - long-term
 - short-term





Schedulers (Cont.)

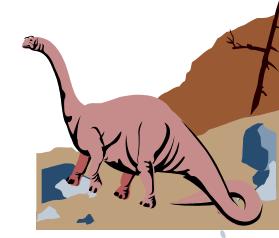
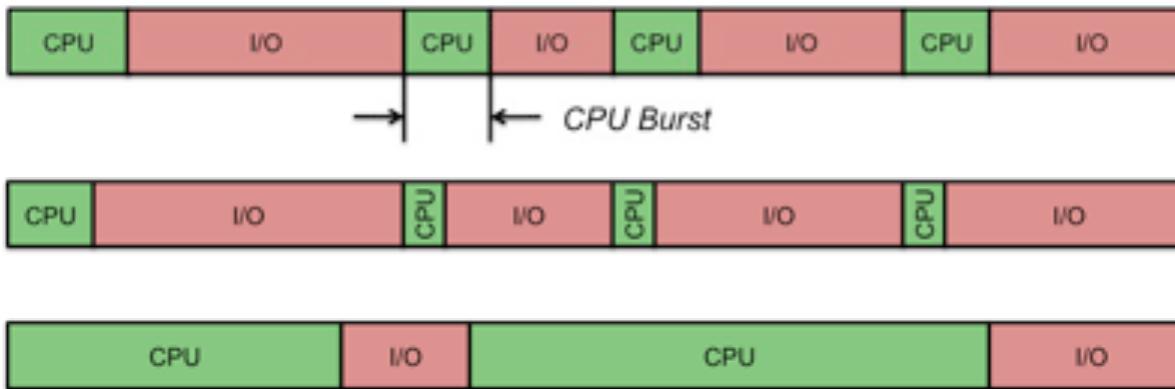
- The long-term scheduler controls the *degree of multiprogramming*.
- Long-term scheduling performs a **gatekeeping function**. It decides whether there's enough memory, or room, to allow new programs into the system.
- Dispatching affects processes
 - ◆ running;
 - ◆ ready;
 - ◆ blocked;
- Long-term scheduling affects processes
 - ◆ new;
 - ◆ exited





Schedulers (Cont.)

- Processes can be described as either:
 - ◆ *I/O-bound process* – spends more time doing I/O than computations, many short CPU bursts.
 - ◆ *CPU-bound process* – spends more time doing computations; few very long CPU bursts.
- The period of computation between I/O requests is called the **CPU burst**.





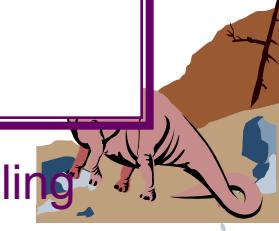
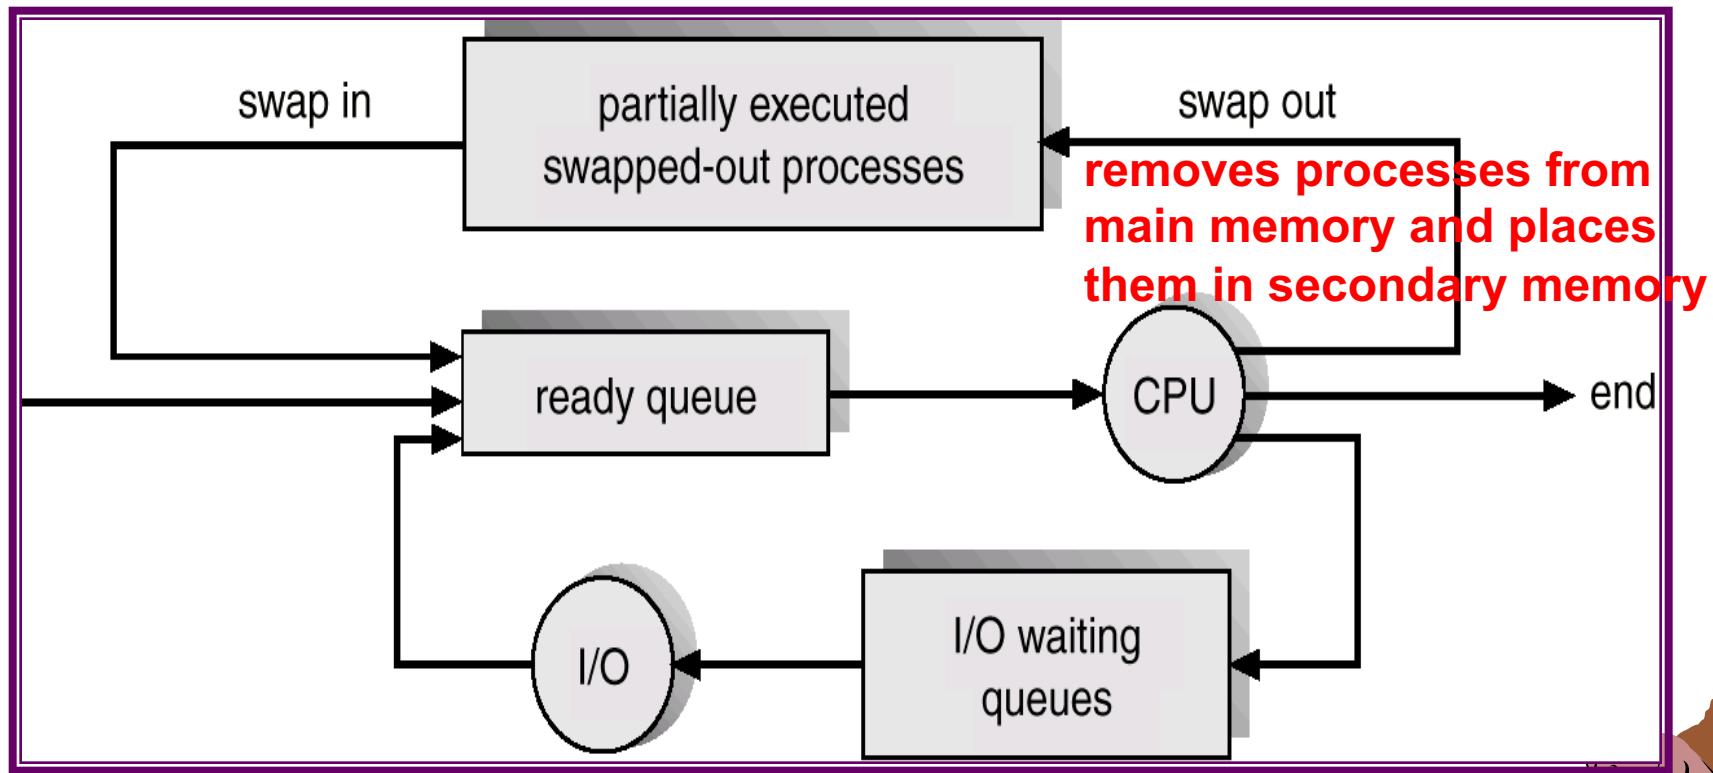
Schedulers (Cont.)

■ Discussion: If you design a CPU scheduler, which type of processes will you give a higher priority of granting CPU resource? I/O-bound processes, or CPU-bound processes?



Addition of Medium-Term Scheduling

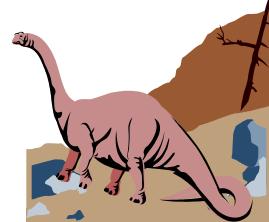
- The resource needs of a process may vary during its runtime. When the system resources become insufficient, some processes may need to swap out





Chapter 3: Processes

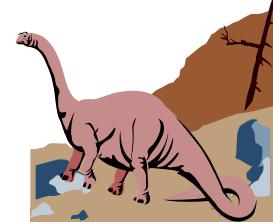
- Process Concept
- Operations and APIs on Processes
- Process Scheduling
- Interprocess Communication
- Communication in Client-Server Systems





Cooperating Processes

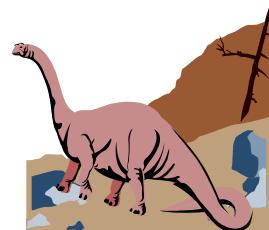
- *Independent* process cannot affect or be affected by the execution of another process.
- *Cooperating* process can affect or be affected by the execution of another process
- Advantages of process cooperation
 - ◆ Information sharing
 - ◆ Computation speed-up
 - ◆ Modularity
 - ◆ Convenience





A Common Cooperating Pattern: Producer-Consumer Problem

- Paradigm for cooperating processes, *producer* process produces information that is consumed by a *consumer* process.
 - ◆ *unbounded-buffer* places no practical limit on the size of the buffer.
 - ◆ *bounded-buffer* assumes that there is a fixed buffer size.





Bounded-Buffer – Share-memory Solution

```
#define BUF_LEN 10
```

```
Typedef struct {
```

```
    ...
```

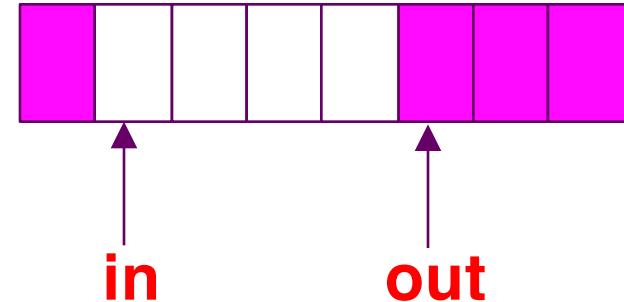
```
} item;
```

```
item buffer[BUF_LEN];
```

```
int in = 0, out = 0;
```

Shared Data

Circular queue



Producer Process

```
item nextProduced;
```

```
while (1) {
```

```
    while (((in+1)%BUF_LEN) == out)
```

```
        /* do nothing */
```

```
    buffer[in] = nextProduced;
```

```
    in = (in + 1) % BUF_LEN;
```

Consumer Process

```
item nextConsumed;
```

```
while (1) {
```

```
    while (in == out)
```

```
        /* do nothing */
```

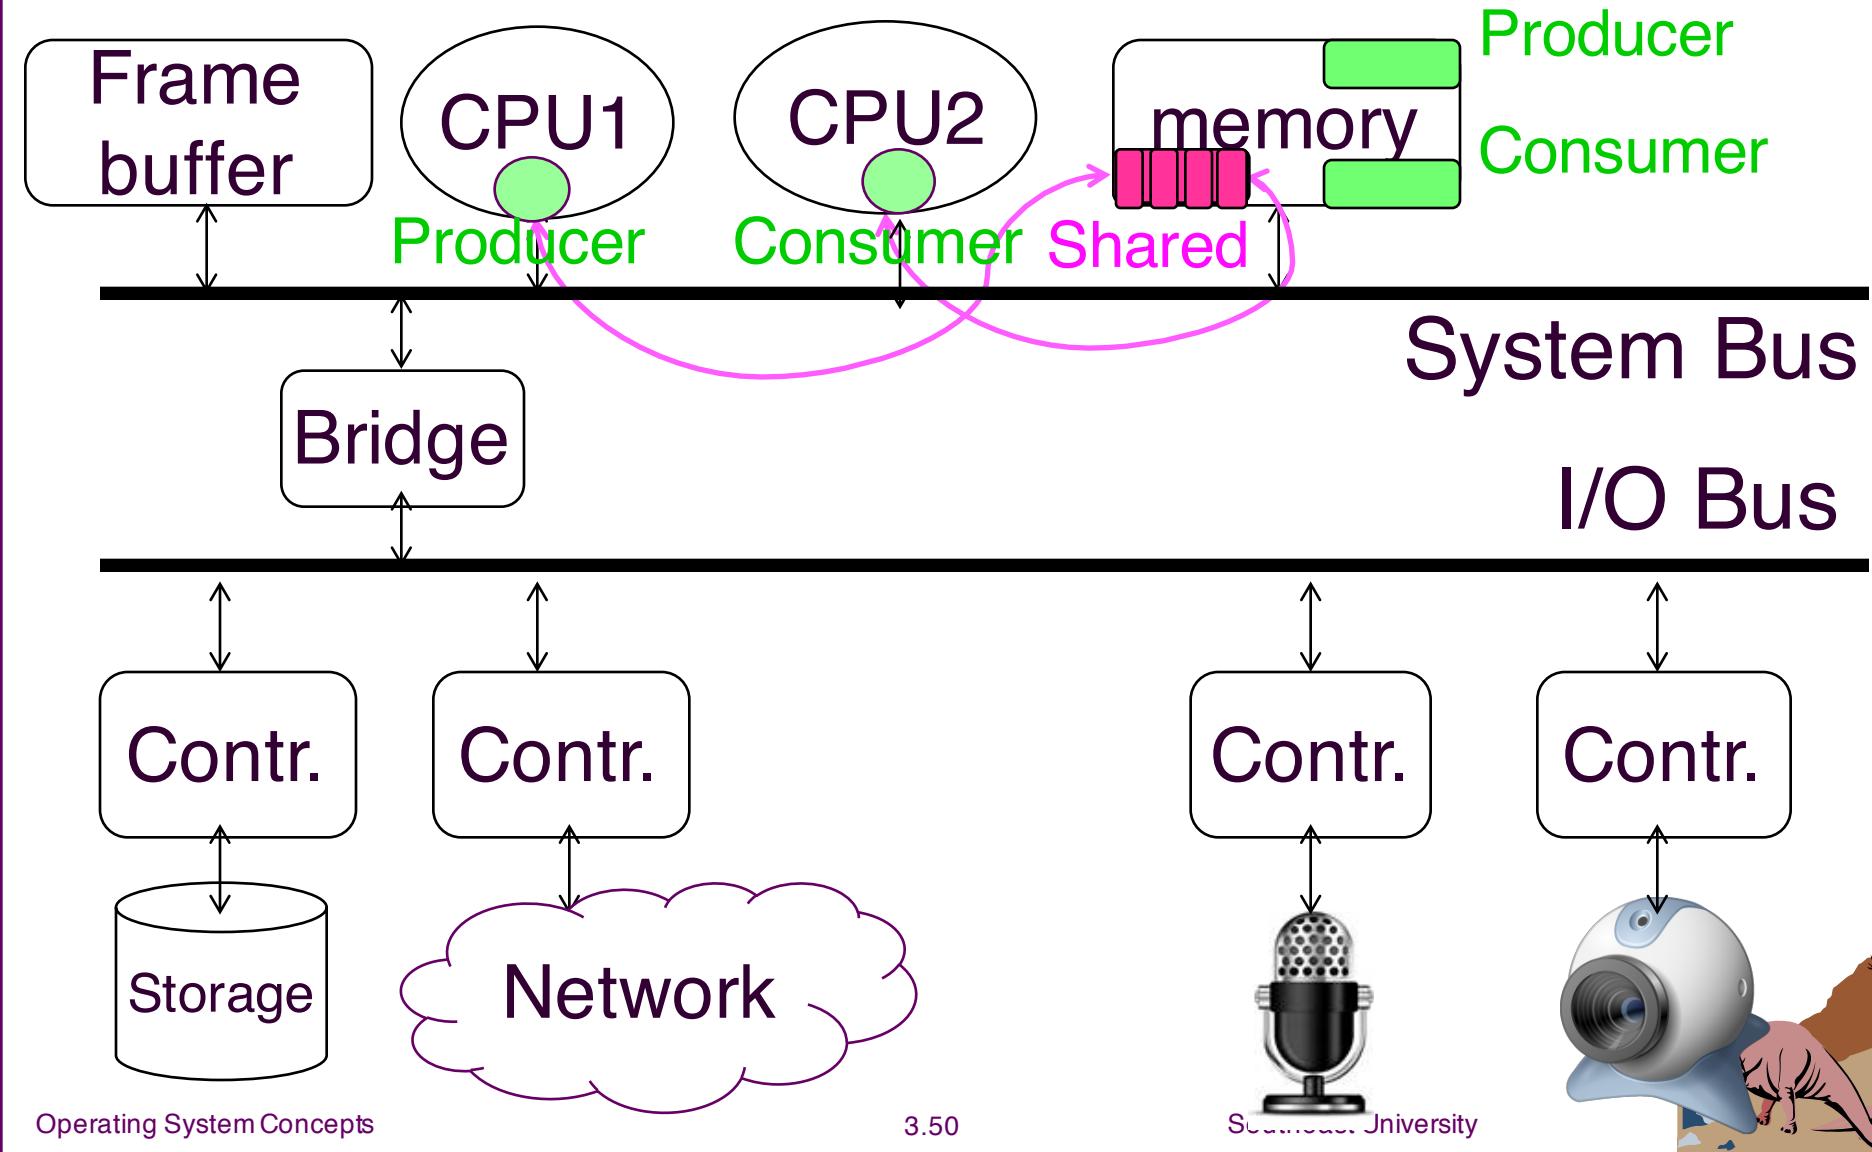
```
    nextConsumed = buffer[out];
```

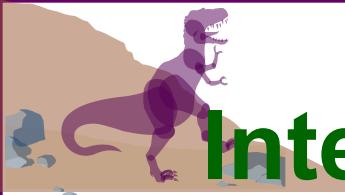
```
    out = (out + 1) % BUF_LEN;
```





Implementation of Communication Link by Shared Memory





Interprocess Communication (IPC)

- Mechanism for processes to communicate and to synchronize their actions.
- Message-passing system – processes communicate with each other without resorting to shared variables.





Interprocess Communication (Cont.)

■ IPC facility provides two operations:

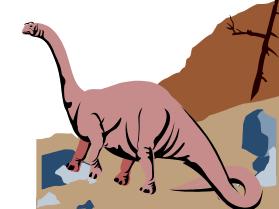
- ◆ **send(*message*)** – message size fixed or variable
- ◆ **receive(*message*)**

■ If P and Q wish to communicate, they need to:

- ◆ establish a *communication link* between them
- ◆ exchange messages via send/receive

■ Implementation of communication link

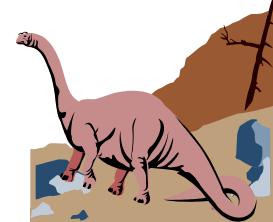
- ◆ physical (e.g., shared memory, hardware bus)
- ◆ logical (e.g., logical properties)





Implementation Questions

- How are links established?
- Can a link be associated with more than two processes?
- How many links can there be between every pair of communicating processes?
- What is the capacity of a link?
- Is the size of a message that the link can accommodate fixed or variable?
- Is a link unidirectional or bi-directional?





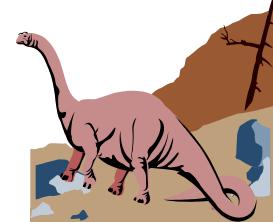
Direct Communication

■ Processes must name each other explicitly:

- ◆ **send** ($P, message$) – send a message to process P
- ◆ **receive** ($Q, message$) – receive a message from process Q

■ Properties of communication link

- ◆ Links are established automatically.
- ◆ A link is associated with exactly one pair of communicating processes.
- ◆ Between each pair there exists exactly one link.
- ◆ The link may be unidirectional, but is usually bi-directional.





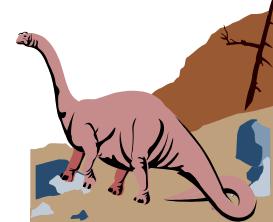
Indirect Communication

■ Messages are directed and received from mailboxes (also referred to as ports).

- ◆ Each mailbox has a unique id.
- ◆ can communicate only if they share a mailbox.

■ Properties of communication link

- ◆ Link established only if processes share a common mailbox
- ◆ A link may be associated with many processes.
- ◆ Each pair of processes may share several communication links.
- ◆ Link may be unidirectional or bi-directional.





Indirect Communication

■ Operations

- ◆ create a new mailbox
- ◆ send and receive messages through mailbox
- ◆ destroy a mailbox

■ Primitives are defined as:

send(A , $message$) – send a message to mailbox A

receive(A , $message$) – receive a message from
mailbox A





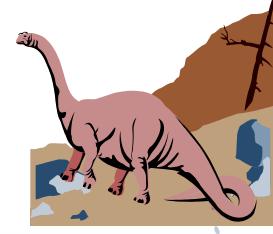
Indirect Communication

■ Mailbox sharing

- ◆ P_1 , P_2 , and P_3 share mailbox A.
- ◆ P_1 , sends; P_2 and P_3 receive.
- ◆ Who gets the message?

■ Solutions

- ◆ Allow a link to be associated with at most two processes.
- ◆ Allow only one process at a time to execute a receive operation.
- ◆ Allow the system to select arbitrarily the receiver.
Sender is notified who the receiver was.





Synchronization

- Message passing may be either blocking or non-blocking.
- **Blocking** is considered **synchronous**
- **Non-blocking** is considered **asynchronous**
- **send** and **receive** primitives may be either blocking or non-blocking.



Buffering

■ Queue of messages attached to the link;
implemented in one of three ways.

1. Zero capacity – 0 messages

 Sender must wait for receiver (rendezvous).

2. Bounded capacity – finite length of n messages

 Sender must wait if link full.

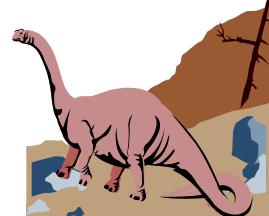
3. Unbounded capacity – infinite length

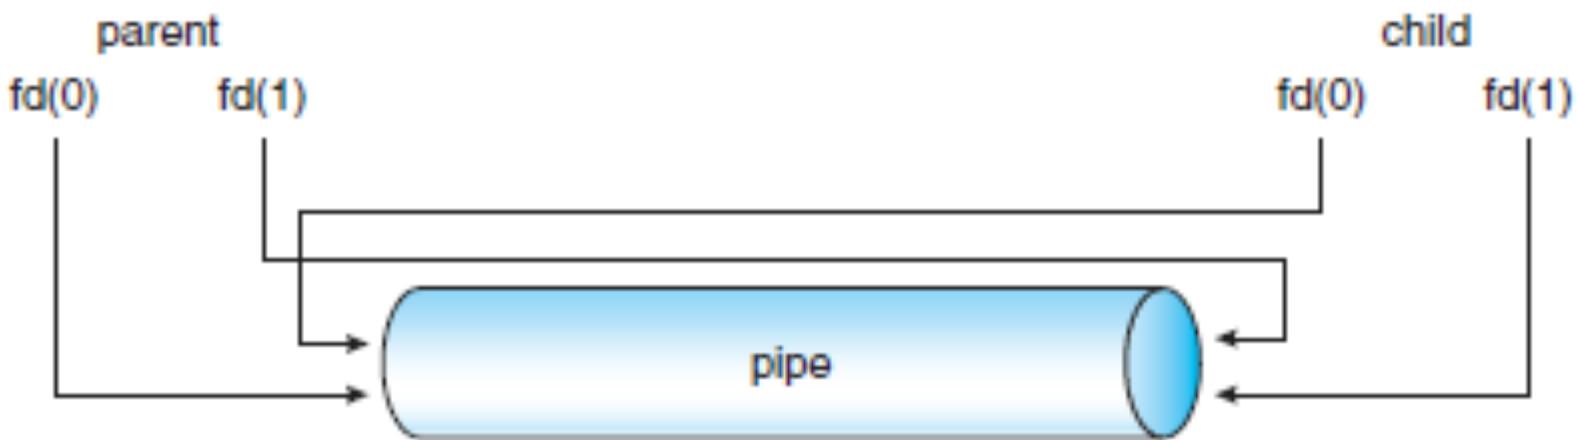
 Sender never blocks.



Pipes in Unix

- UNIX pipes are implemented in a similar way, but with the `pipe()` system call.
 - ◆ The output of one process is connected to an in-kernel pipe.
 - ◆ The input of another process is connected to that same pipe.
 - ◆ E.g.,
 - ✓ `ls | wc`





```

if (pid > 0) { /* parent process */
    /* close the unused end of the pipe */
    close(fd[READ_END]);

    /* write to the pipe */
    write(fd[WRITE_END], write_msg, strlen(write_msg)+1);

    /* close the write end of the pipe */
    close(fd[WRITE_END]);
}

else { /* child process */
    /* close the unused end of the pipe */
    close(fd[WRITE_END]);

    /* read from the pipe */
    read(fd[READ_END], read_msg, BUFFER_SIZE);
    printf("read %s", read_msg);

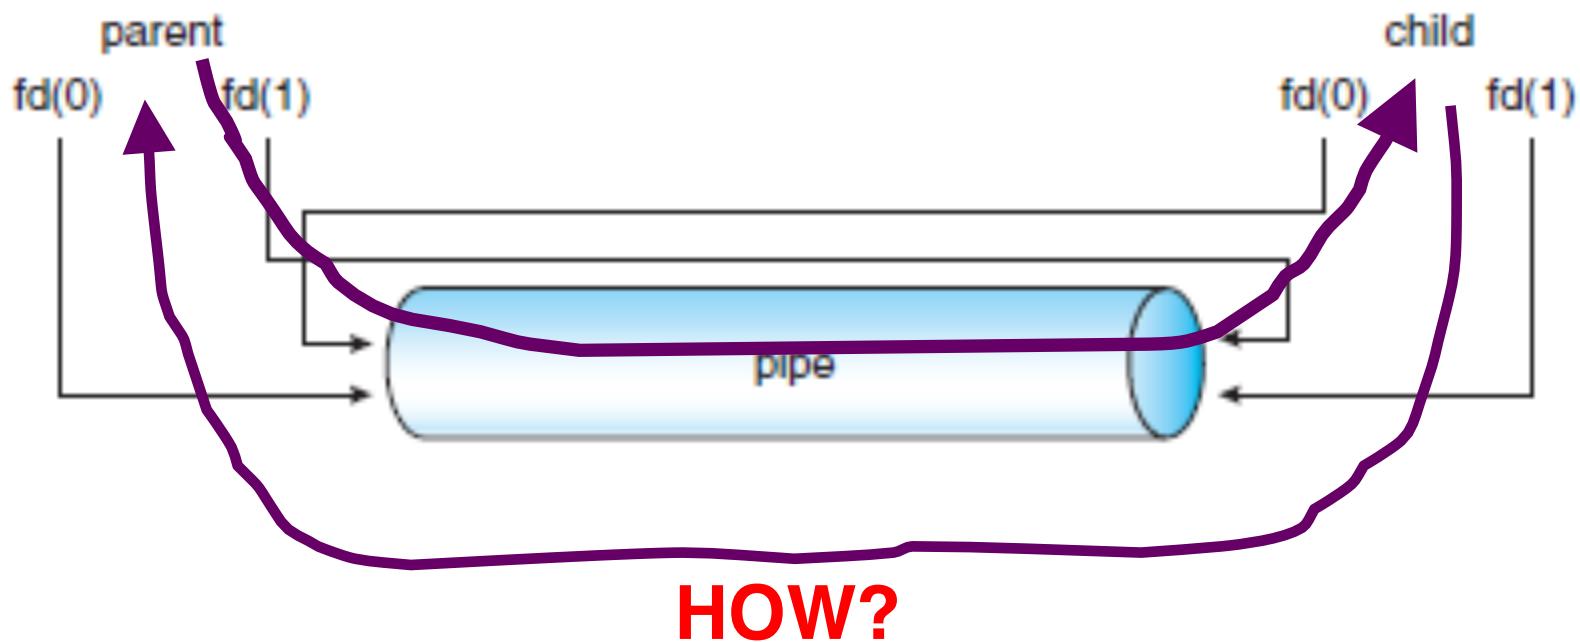
    /* close the write end of the pipe */
    close(fd[READ_END]);
}

```

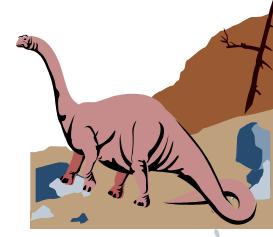


Discussion

- What if the parent wants to write something to child, while child also wants to write something to parent?



- Hints, ordinary pipes are unidirectional





Chapter 3: Processes

- Process Concept
- Operations and APIs on Processes
- Process Scheduling
- Cooperating Processes
- Interprocess Communication
- Communication in Client-Server Systems



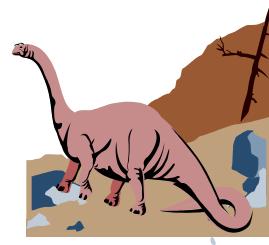


Client-Server Communication

■ Sockets

■ Remote Procedure Calls

■ Remote Method Invocation (Java)



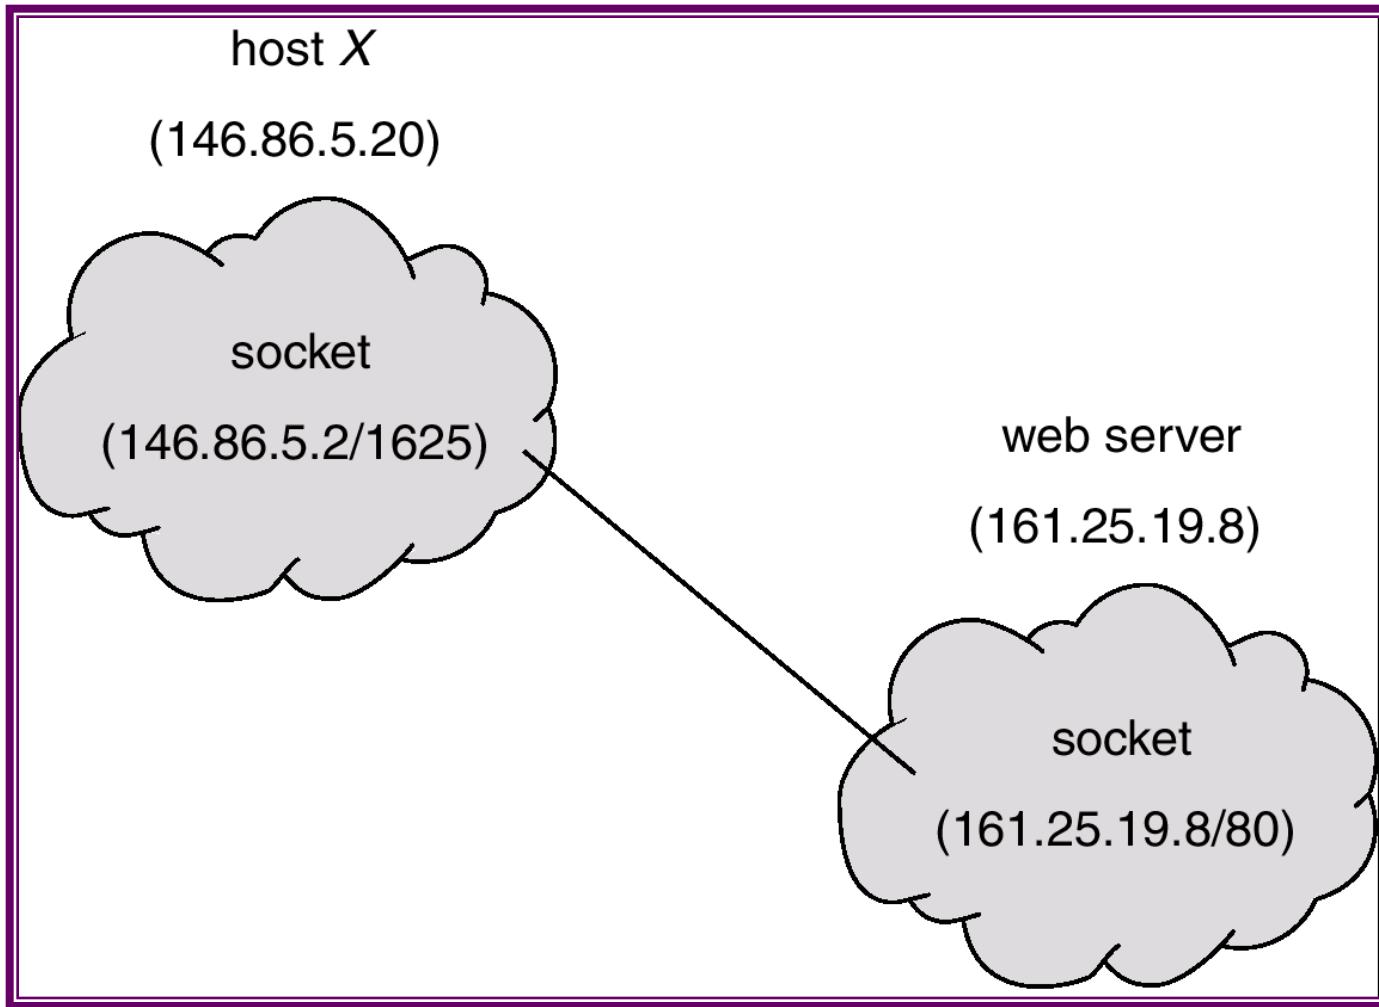


Sockets

- A socket is defined as an *endpoint for communication*.
- Concatenation of IP address and port
- The socket **161.25.19.8:1625** refers to port **1625** on host **161.25.19.8**
- Communication consists between a pair of sockets.



Socket Communication



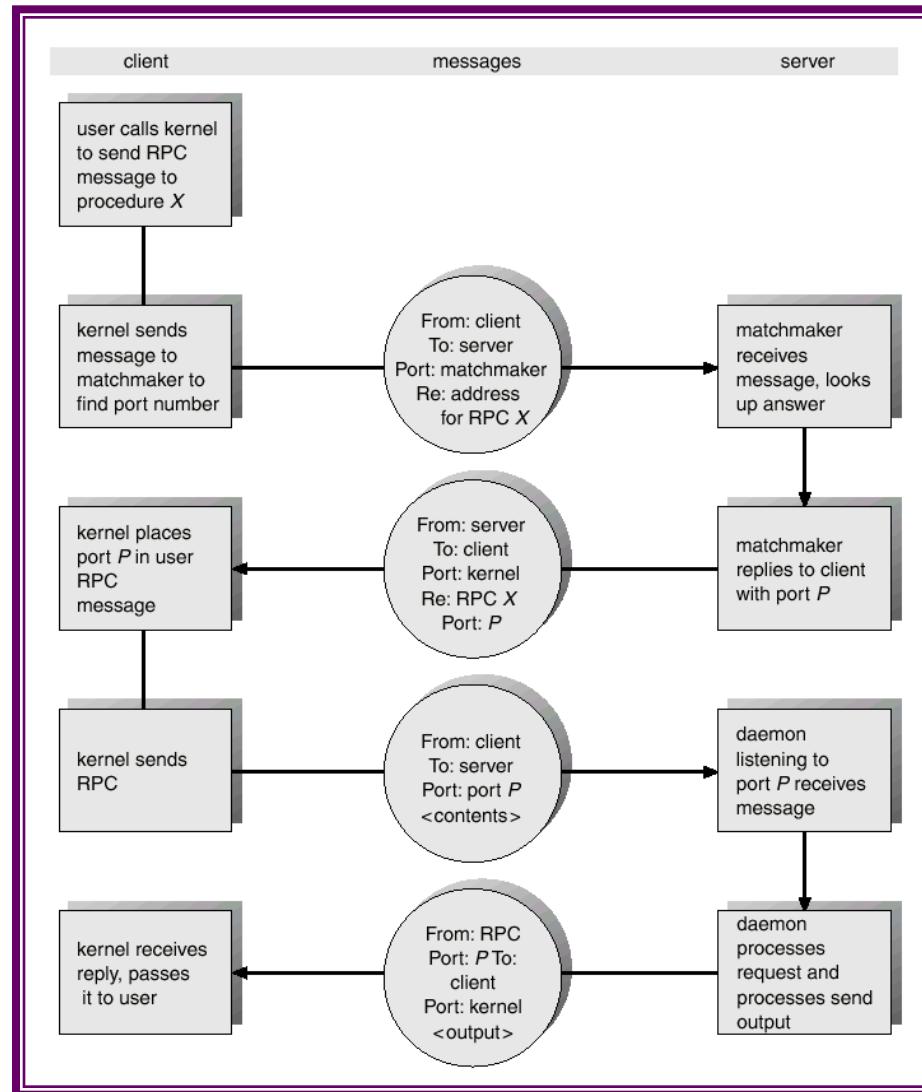


Remote Procedure Calls

- Remote procedure call (RPC) abstracts procedure calls between processes on networked systems.
- **Stubs** – client-side proxy for the actual procedure on the server.
- The client-side stub locates the server and *marshalls* the parameters.
- The server-side stub receives this message, unpacks the marshalled parameters, and performs the procedure on the server.



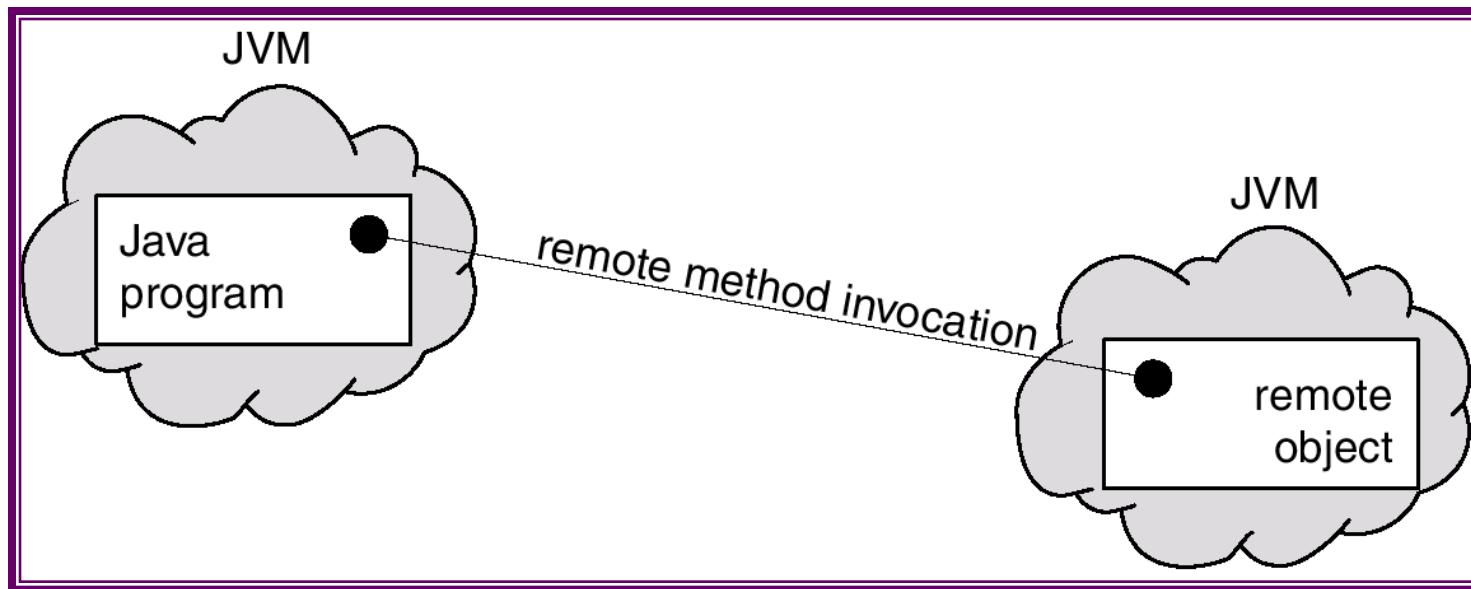
Execution of RPC





Remote Method Invocation

- Remote Method Invocation (RMI) is a Java mechanism similar to RPCs.
- RMI allows a Java program on one machine to invoke a method on a remote object.





Marshalling Parameters

