

Chapter 8: Memory Management

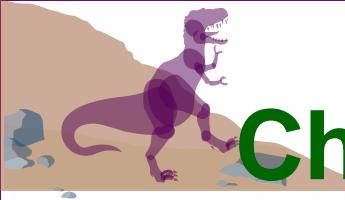
肖卿俊

办公室：九龙湖校区计算机楼212室

电邮：csqjxiao@seu.edu.cn

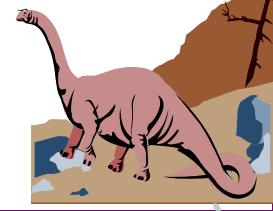
主页：<https://csqjxiao.github.io/PersonalPage>

电话：025-52091022



Chapter 8: Memory Management

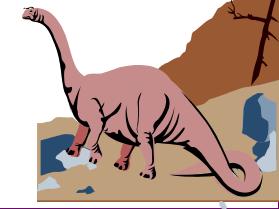
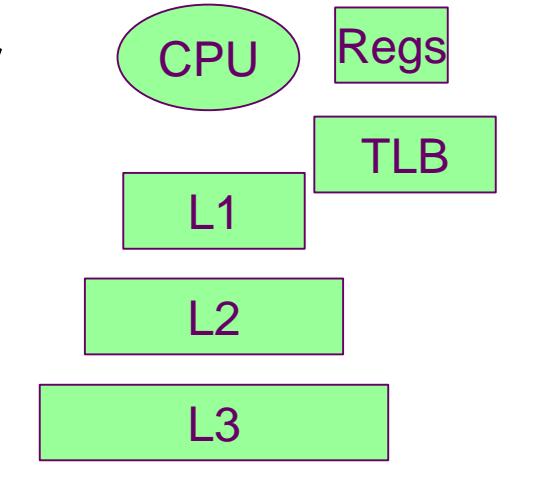
- Background
- Swapping
- Contiguous Allocation
- Segmentation
- Paging
- Advanced Page Table Structure
- Segmentation with Paging

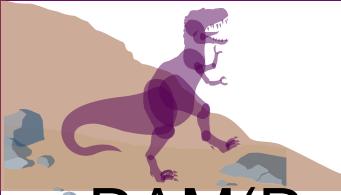




Background for Memory Hierarchy

- Main memory and registers are the only storage CPU can access directly
- Register access in one CPU clock (or less)
- Main memory can take many cycles
- **L1/L2/L3 Cache** sits between main memory and CPU registers
- Protection of memory required to ensure correct operation
- Program must be brought into main memory and placed within a process for it to be run.





什么是RAM?

■ RAM(Random Access Memory) 的中文是随机存取存储器

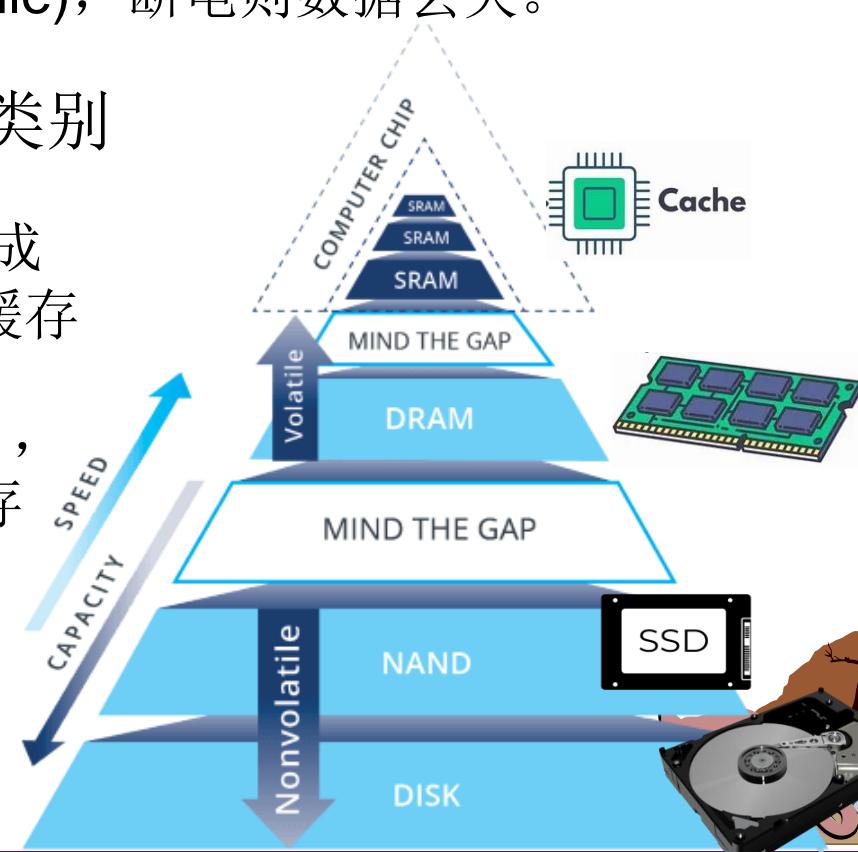
- 为什么要强调随机存储呢？因为在此之前，大部分的存储器都是顺序存储（**Direct-Access**），比较常见的如硬盘，光碟，老式的磁带，磁鼓存储器等等。随机存取存储器的特点是其访问数据的时间与数据存放在存储器中的物理位置无关。
- RAM的另一个特点是易失性（**Volatile**），断电则数据丢失。

■ RAM主要分SRAM和DRAM两种类别

- SRAM是静态RAM (**Static RAM**)，成本高，存取速度比较快，用于高速缓存

- DRAM是动态RAM (**Dynamic RAM**)，成本低，存取速度比较慢，用于内存

■ 本章的内存管理主要是DRAM资源在进程内部和之间的分配算法





OS Commands to get CPU and memory Information

■ Linux

- ◆ **cat /proc/cpuinfo**
- ◆ **cat /proc/meminfo**

```
% system_profiler SPHardwareDataType
```

Hardware Overview:

Model Name: MacBook Pro

Model Identifier: MacBookPro15,1

Processor Name: 6-Core Intel Core i7

Processor Speed: 2.2 GHz

Number of Processors: 1

Total Number of Cores: 6

L2 Cache (per Core): 256 KB

L3 Cache: 9 MB

Hyper-Threading Technology: Enabled

Memory: 16 GB

System Firmware Version: 1731.100.130.0.0

OS Loader Version: 540.100.7~23

Serial Number (system): C02XH3U4JG5L

Hardware UUID: AE627AB2-1D36-5908-BCB5-AE7E7E83D110

Provisioning UDID: AE627AB2-1D36-5908-BCB5-AE7E7E83D110

85

■ MacOS

- **system_profiler SPHardwareDataType**
- **vm_stat**

```
% vm_stat
```

Mach Virtual Memory Statistics:

(page size of 4096 bytes)

Pages free: 96730.

Pages active: 1428706.

Pages inactive: 1262257.

Pages speculative: 168869.

Pages throttled: 0.

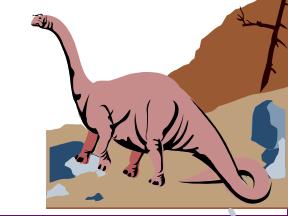
Pages wired down: 987490.

Pages purgeable: 84220.

"Translation faults":3533751188.

Pages copy-on-write:396805612.

Pages zero filled: 3188592824.



Reminder: A Typical Memory Hierarchy

- Everything is a cache for something else

	Access time	Capacity	Managed By
On the datapath	Registers	1 cycle	1 KB Software/Compiler
On chip	Level 1 Cache	2-4 cycles	32 KB Hardware
	Level 2 Cache	10 cycles	256 KB Hardware
	Level 3 Cache	40 cycles	10 MB Hardware
Other chips	Main Memory	200 cycles	10 GB Software/OS
Mechanical devices	Flash Drive	10-100us	100 GB Software/OS
	Hard Disk	10ms	1 TB Software/OS


The diagram illustrates a memory hierarchy with the following components and their relative positions:

- Registers:** Labeled "On the datapath".
- Level 1 Cache:** Labeled "On chip".
- Level 2 Cache:** Labeled "On chip".
- Level 3 Cache:** Labeled "On chip".
- Main Memory:** Labeled "Other chips".
- Flash Drive:** Labeled "Mechanical devices".
- Hard Disk:** Labeled "Mechanical devices".


Vertical dashed arrows indicate the data path between adjacent components in the hierarchy.




Background on Multistep Processing of a User Program



Compile time is the period when the programming code is converted to machine code.



Load time is the duration it takes for a computer program to be loaded into memory and become ready for execution.





Binding of Instructions and Data to Physical Memory Addresses

■ Compile time

- ◆ If memory location of running a program is known a priori, absolute code can be generated by compiler; must recompile code if starting location changes.

■ Load time

- ◆ Must generate *relocatable* code if memory location is not known at compile time

■ Run time

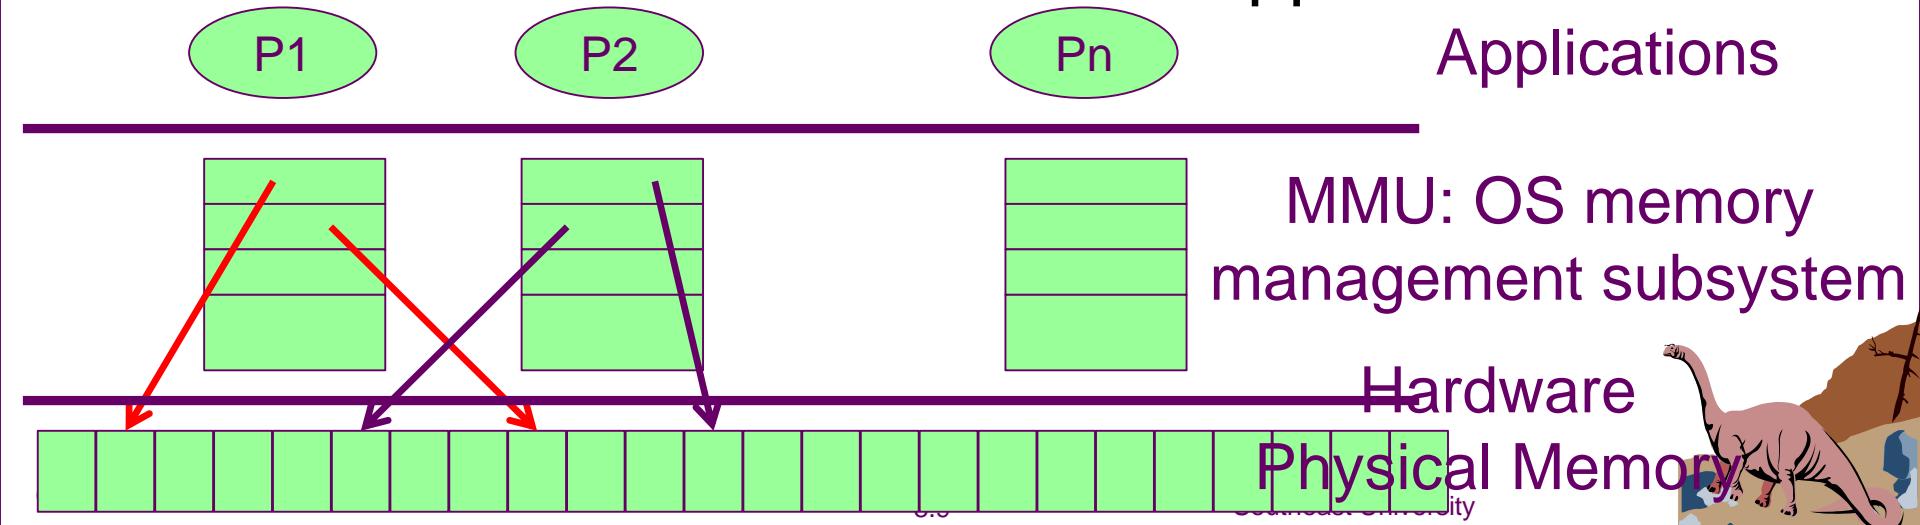
Most general-purpose operating systems use the execution-time address binding

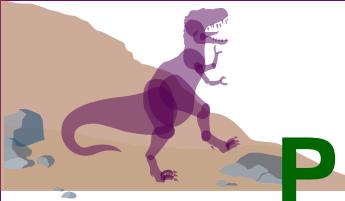
- ◆ Binding delayed until run time if the process can be moved during its execution from one memory segment to another. Need hardware support for address mappings.



Logical Address Space vs. Physical Address Space

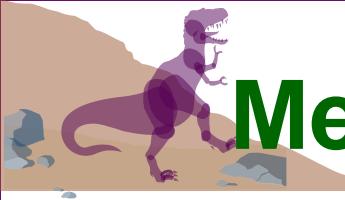
- The concept of a *logical address space* that is bound to a separate *physical address space* is central to the proper memory management.
 - ◆ *Logical address* – generated by the CPU; also referred to as *virtual address*.
 - ◆ *Physical address* – address seen by main memory units. LOGI addr are runtime mapped to PHY addr.






Logical Address Space vs. Physical Address Space (cont.)

- Logical and physical addresses are the same in **compile-time** and **load-time** address-binding schemes
- Logical and physical addresses differ in **execution-time address-binding** scheme.
- In this case, logical address is also referred to as virtual address. (Logical = Virtual in this course)



Memory-Management Unit (MMU)



- **Hardware device** that maps virtual to physical address.
- In MMU scheme, the value in the **relocation register** is added to every address generated by a user process at the time it is sent to memory.
- The user program deals with *logical* addresses; it never sees the *real* physical addresses.



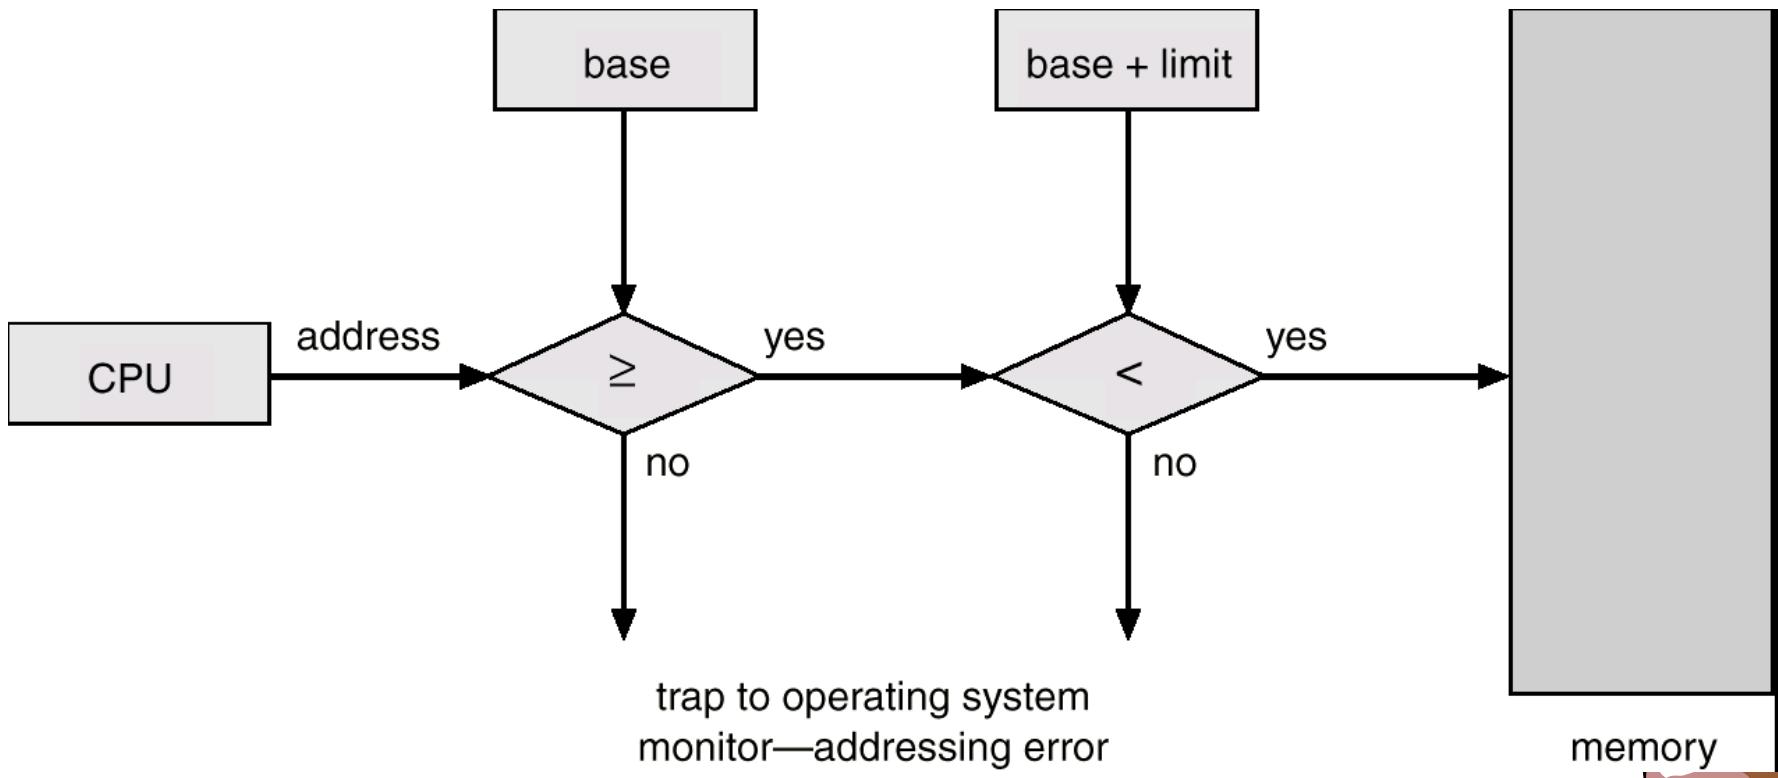


Revisit the Simple Memory Management: Base + Limit Registers

- A pair of **base** and **limit** registers can define the address space of each process



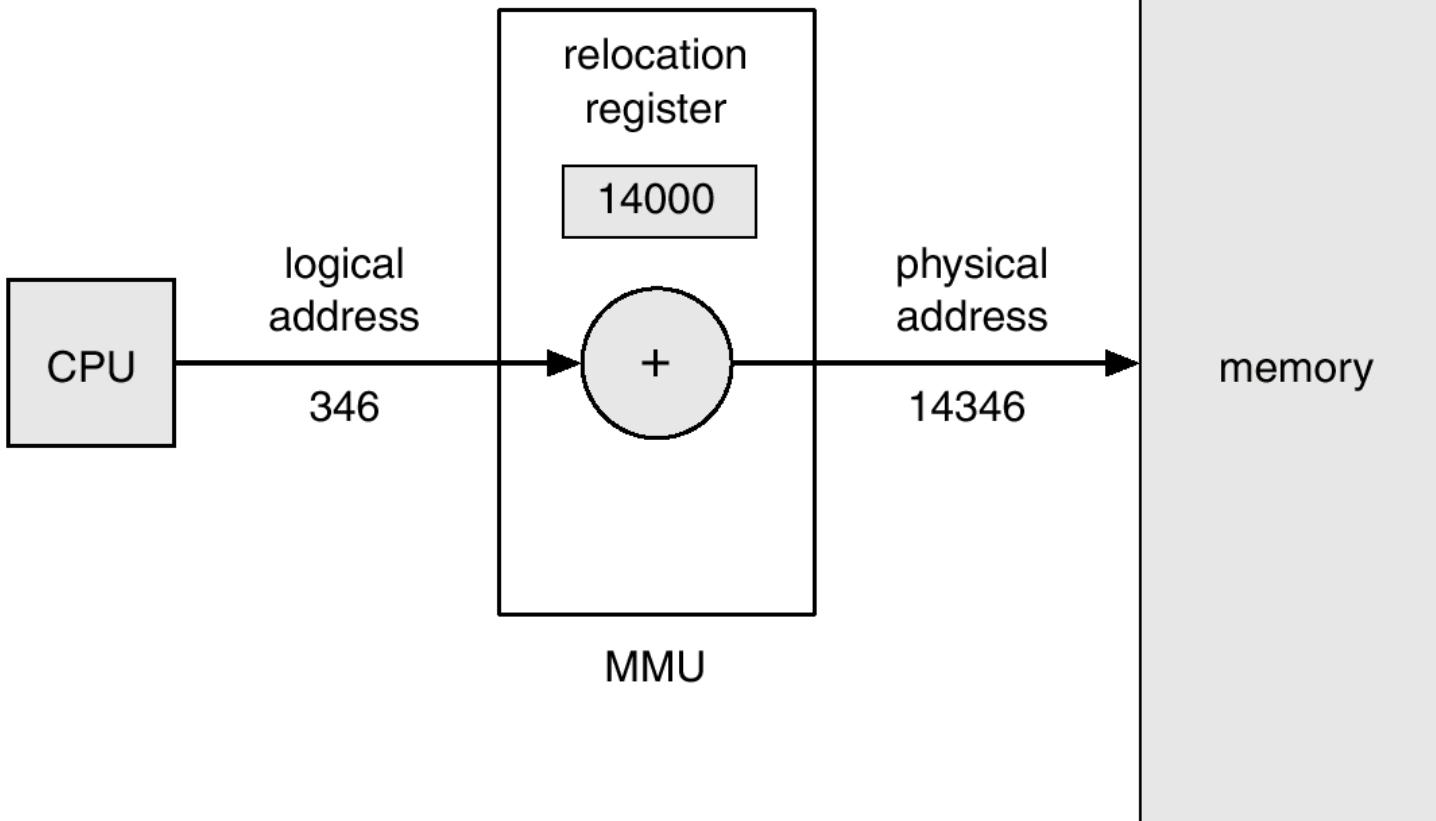
Applications access memory units by physical addresses, when there is no separation between virtual and physical addresses





However, applications only know logical addresses. So use a **relocation register** instead of a base register

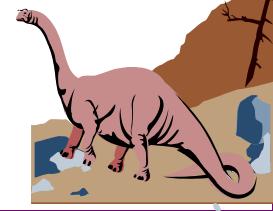
A Simple MMU





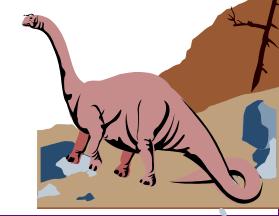
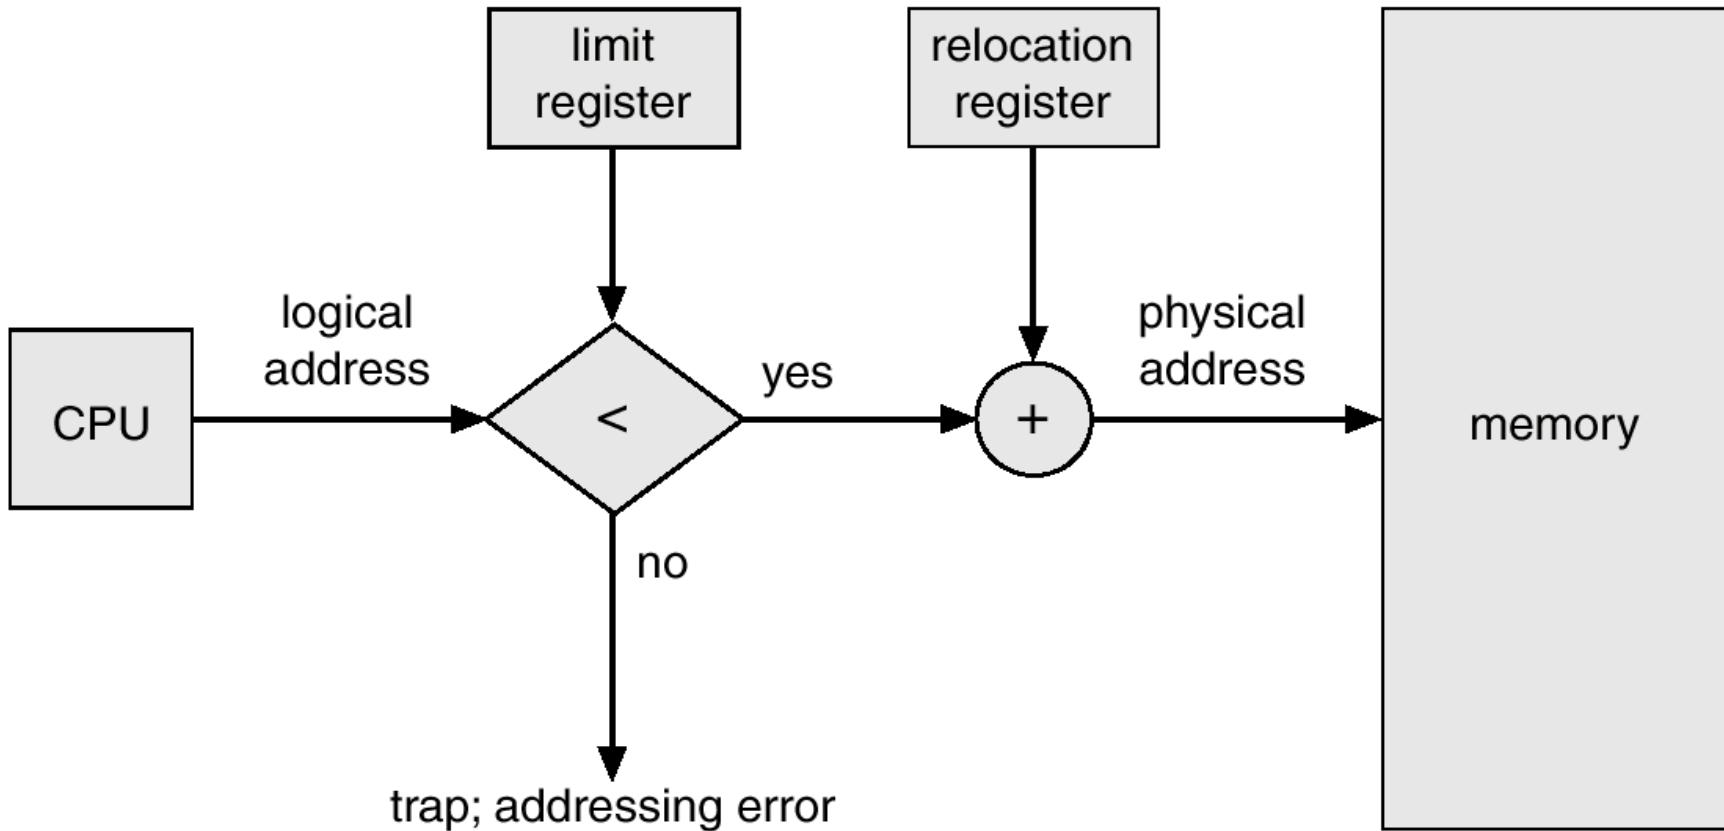
Memory Protection

- Relocation-register scheme used to protect user processes from each other, and from changing operating-system code and data.
- Relocation register contains value of the smallest physical address
- Limit register contains range of logical addresses – each logical address must be less than the limit register.





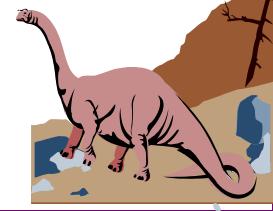
Hardware Support for Relocation and Limit Registers





Chapter 9: Memory Management

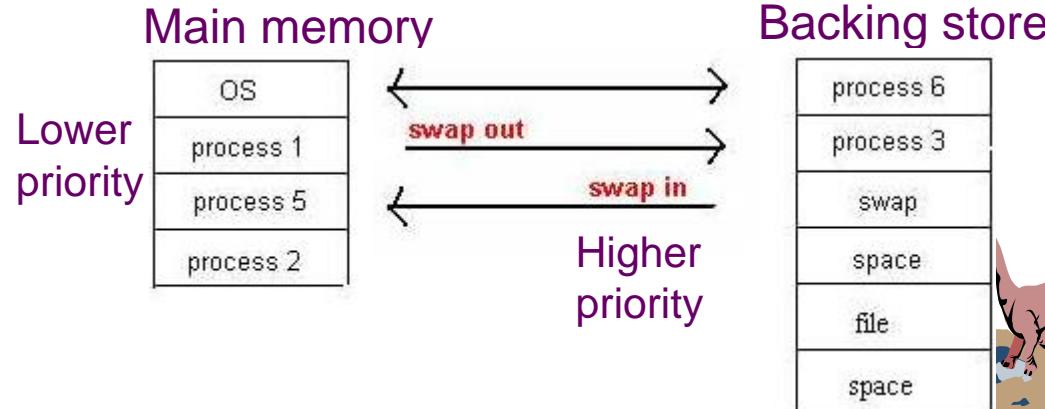
- Background
- Swapping
- Contiguous Allocation
- Segmentation
- Paging
- Advanced Page Table Structure
- Segmentation with Paging





Schematic View of Swapping

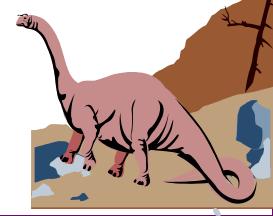
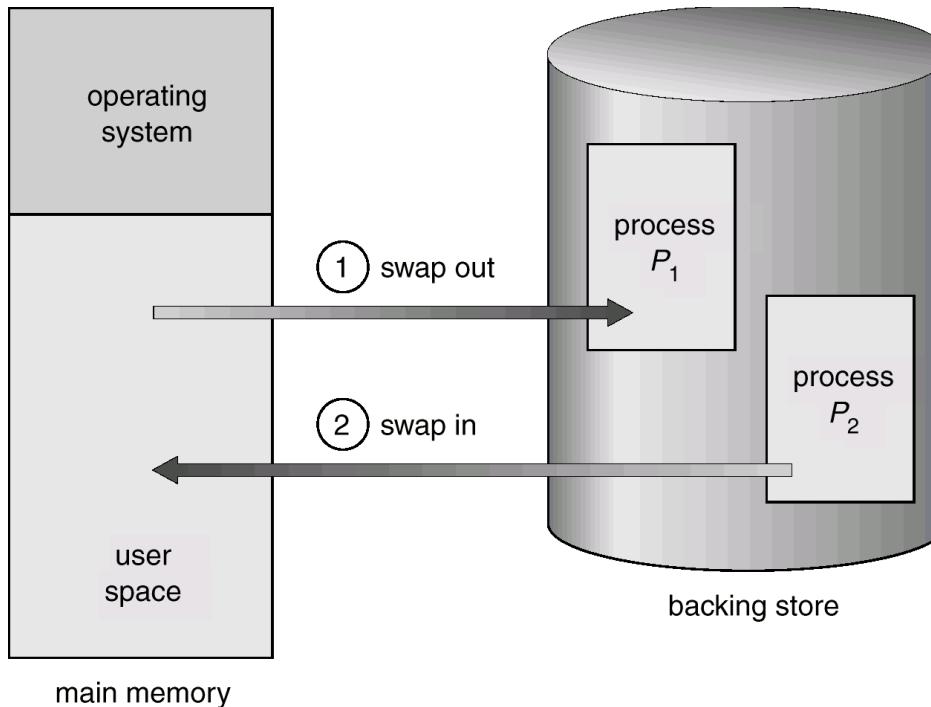
- A process can be swapped temporarily out of memory to a *Backing store*, and then brought back into memory for continued execution.
- Backing store – fast disk large enough to hold copies of all memory images for all users; must provide direct access to these memory images.
- *Roll out, roll in* – swapping variant used for priority-based scheduling algorithms; lower-priority process is swapped out so higher-priority process can be loaded and executed.





Major Time Overhead of Swapping


- Main memory reference 100ns
- Magnetic Disk track seek 10,000,000 ns
- Major part of swap time is transfer time; total transfer time is directly proportional to the amount of memory swapped.





Chapter 9: Memory Management

- Background
- Swapping
- Contiguous Allocation
- Segmentation
- Paging
- Advanced Page Table Structure
- Segmentation with Paging

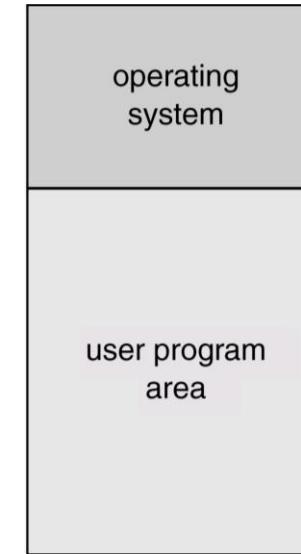




Contiguous Allocation


■ Monoprogramming systems usually have two partitions:

- ◆ Resident operating system, usually held in low memory with interrupt vector.
- ◆ User processes then held in high memory.



■ Multiprogramming Systems:

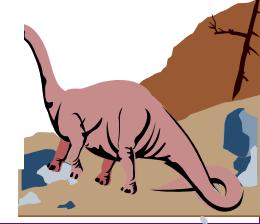
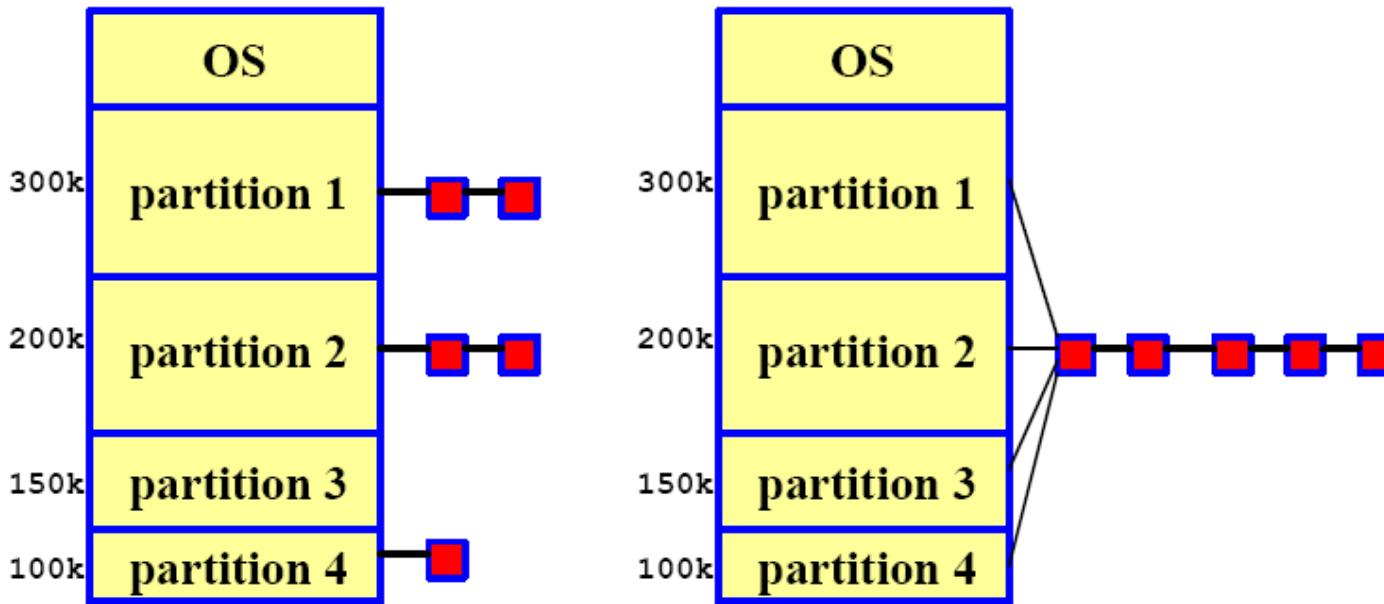
- ◆ Fixed partitions
- ◆ Variable partitions





Fixed Partitions


- Main memory is divided into n partitions.
- Partitioning can be done at the startup time and altered later on.
- Each partition may have a job queue. Or, all partitions share the same job queue.





Variable Partitions

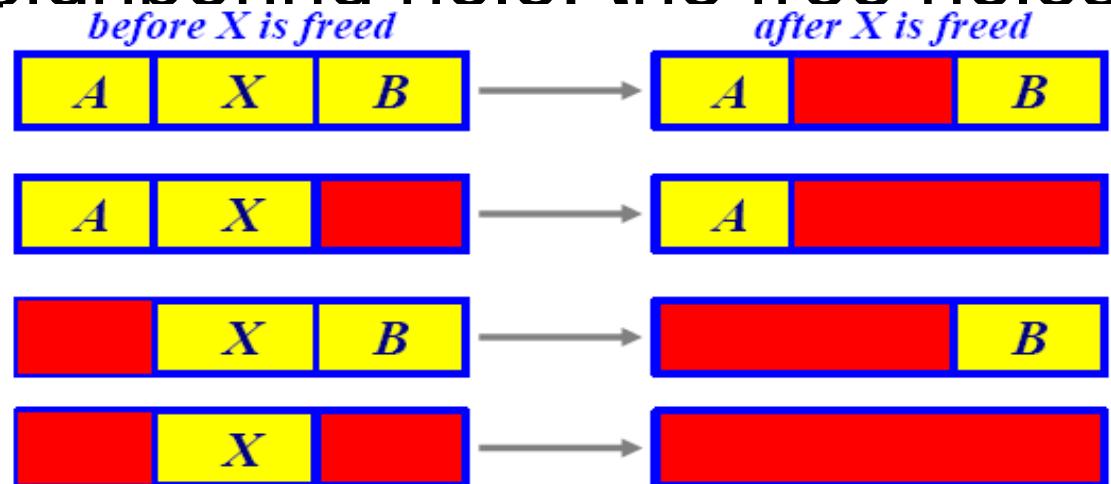
- **Hole** – block of available memory; holes of various size are scattered throughout memory.
- When a process arrives, it is allocated memory from a hole large enough to accommodate it.
- Thus, partition sizes are not fixed, The number of partitions also varies.
- Operating system maintains information about:
 - allocated partitions
 - free partitions(hole)





List of Free Holes

- If the hole is larger than the requested size, it is cut into two. The one of the requested size is given to the process, the remaining one becomes a *new* hole.
- When a process returns a memory block, it becomes a hole and must be merged with its neighbor.
- For finding the neiahborina hole. the free holes are organized as *addresses*.






Dynamic Storage-Allocation Problem

How to satisfy a request of size n from a list of free holes.

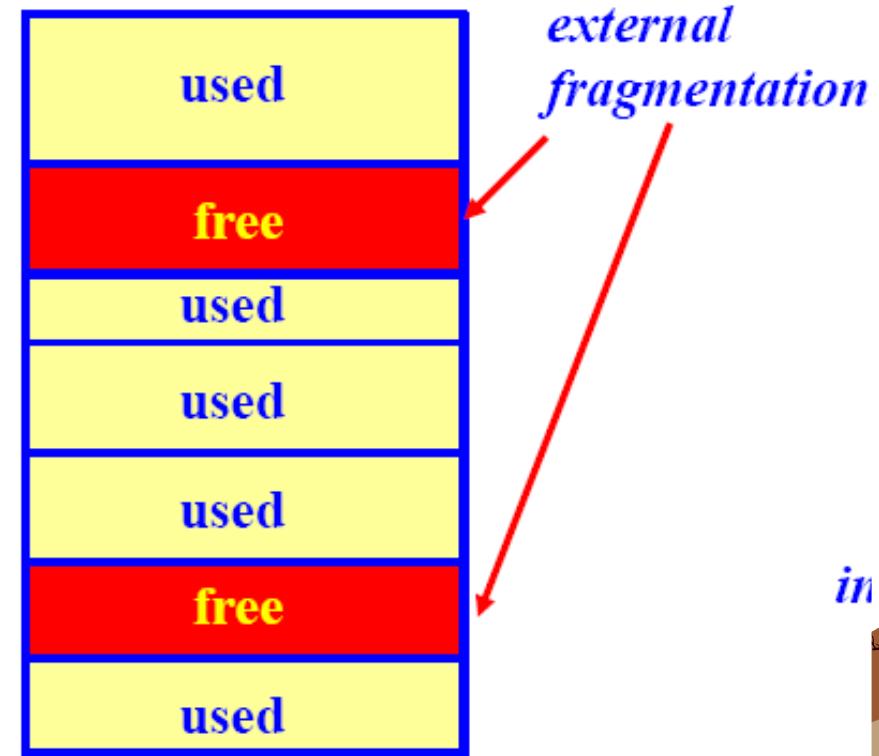
- **First-fit (首次适配)**: Allocate the *first* hole that is big enough.
- **Best-fit (最佳适配)**: Allocate the *smallest* hole that is big enough; must search entire list.
Produces the smallest leftover hole.
- **Worst-fit (最差适配)**: Allocate the *largest* hole; must also search entire list. Produces the largest leftover hole.





External Fragmentation (外部内存碎片)


- Processes are loaded and removed from memory. Eventually, the memory will be cut into small holes that are not large enough to run any incoming process.
- Free memory holes between allocated ones are called *external fragmentation*.





Internal Fragmentation (内部内存碎片)

- It is unwise to allocate exactly the requested amount of memory to a process, because of the minimum requirement for memory management.
- Thus, memory that is allocated to a partition, but is not used, are called *internal fragmentation*.

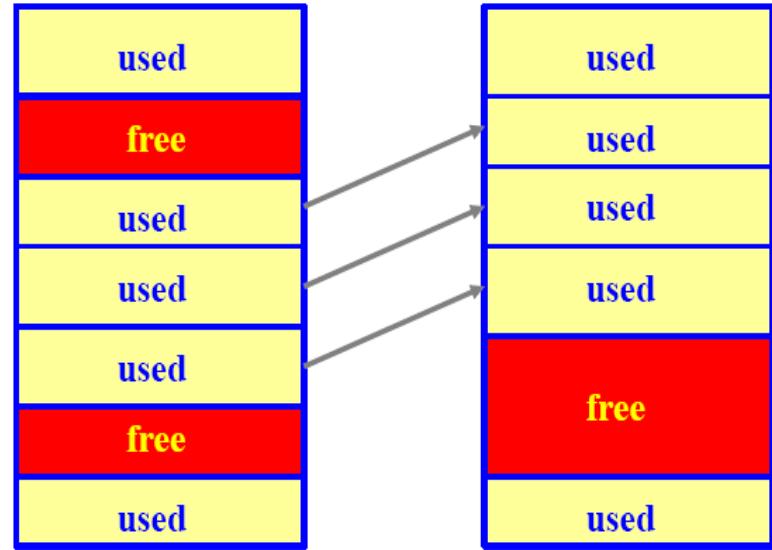




Compaction for Less External Fragmentation

- Shuffle memory contents to place all free memory together in one large block.
- Compaction is possible *only* if program relocation is dynamic and is done at execution time.
- Runtime compaction scheme can be time-consuming

Runtime Memory Compaction

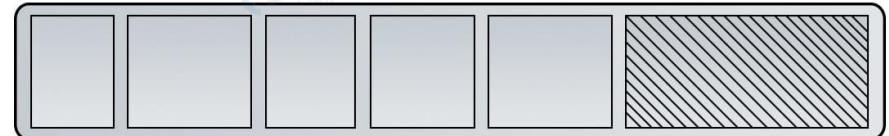


OR illustrated as

Fragmented Memory Before Compaction



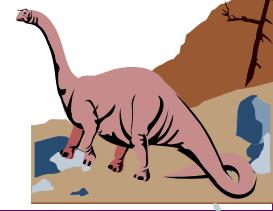
Memory After Compaction





Chapter 9: Memory Management

- Background
- Swapping
- Contiguous Allocation
- Segmentation
- Paging
- Advanced Page Table Structure
- Segmentation with Paging






Segmentation

- Memory-management scheme that supports **user view of memory**.
- A program is a collection of segments.
- A segment is a logical unit such as:
 - main program,
 - procedure or function,
 - method,
 - object,
 - local variables,
 - global variables,
 - common block,
 - stack,
 - symbol table, arrays

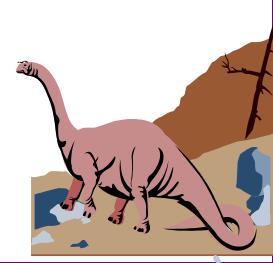
User's View of a Program





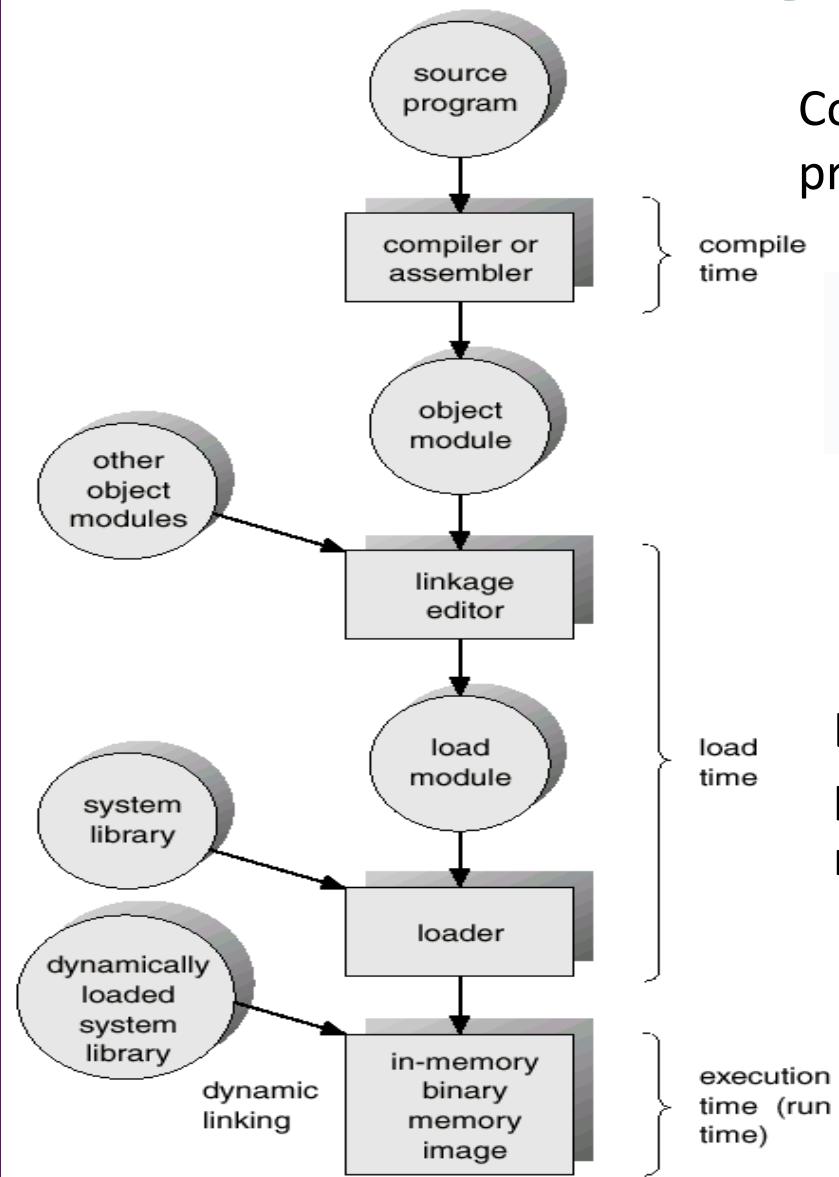
A Previously Used Example

```
//main.cpp
int a = 0; ← 数据段, 全局变量
char *p1; ← 数据段, 全局变量
main()
{
    int b; ← 栈段, 局部变量
    char s[] = "abc"; ← 栈段, 局部变量
    char *p2; ← 栈段, 局部变量
    char *p3 = "123456"; ← 栈段, 局部变量
    p1 = (char *)malloc(10); ← 堆段
    p2 = (char *)malloc(20); ← 堆段
}
```






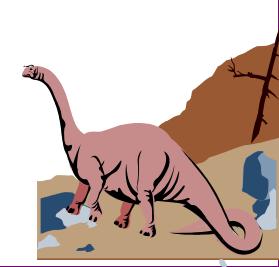
Background on Multistep Processing of a User Program

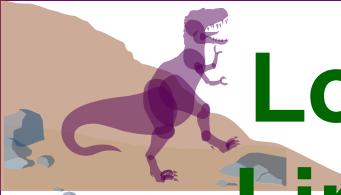


Compile time is the period when the programming code is converted to machine code.

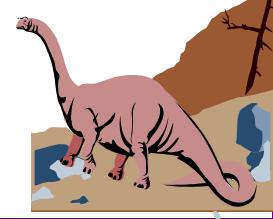



Load time is the duration it takes for a computer program to be loaded into memory and become ready for execution.







Loading an ELF (Executable and Linkable Format) binary on Linux






Mapping Segments to Physical Memory

- Each segment occupies a contiguous memory space in the physical memory



physical memory space

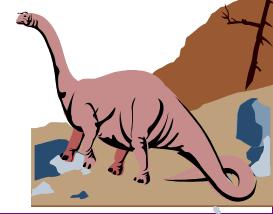




Segmentation Architecture

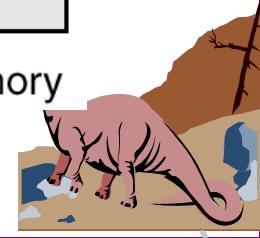
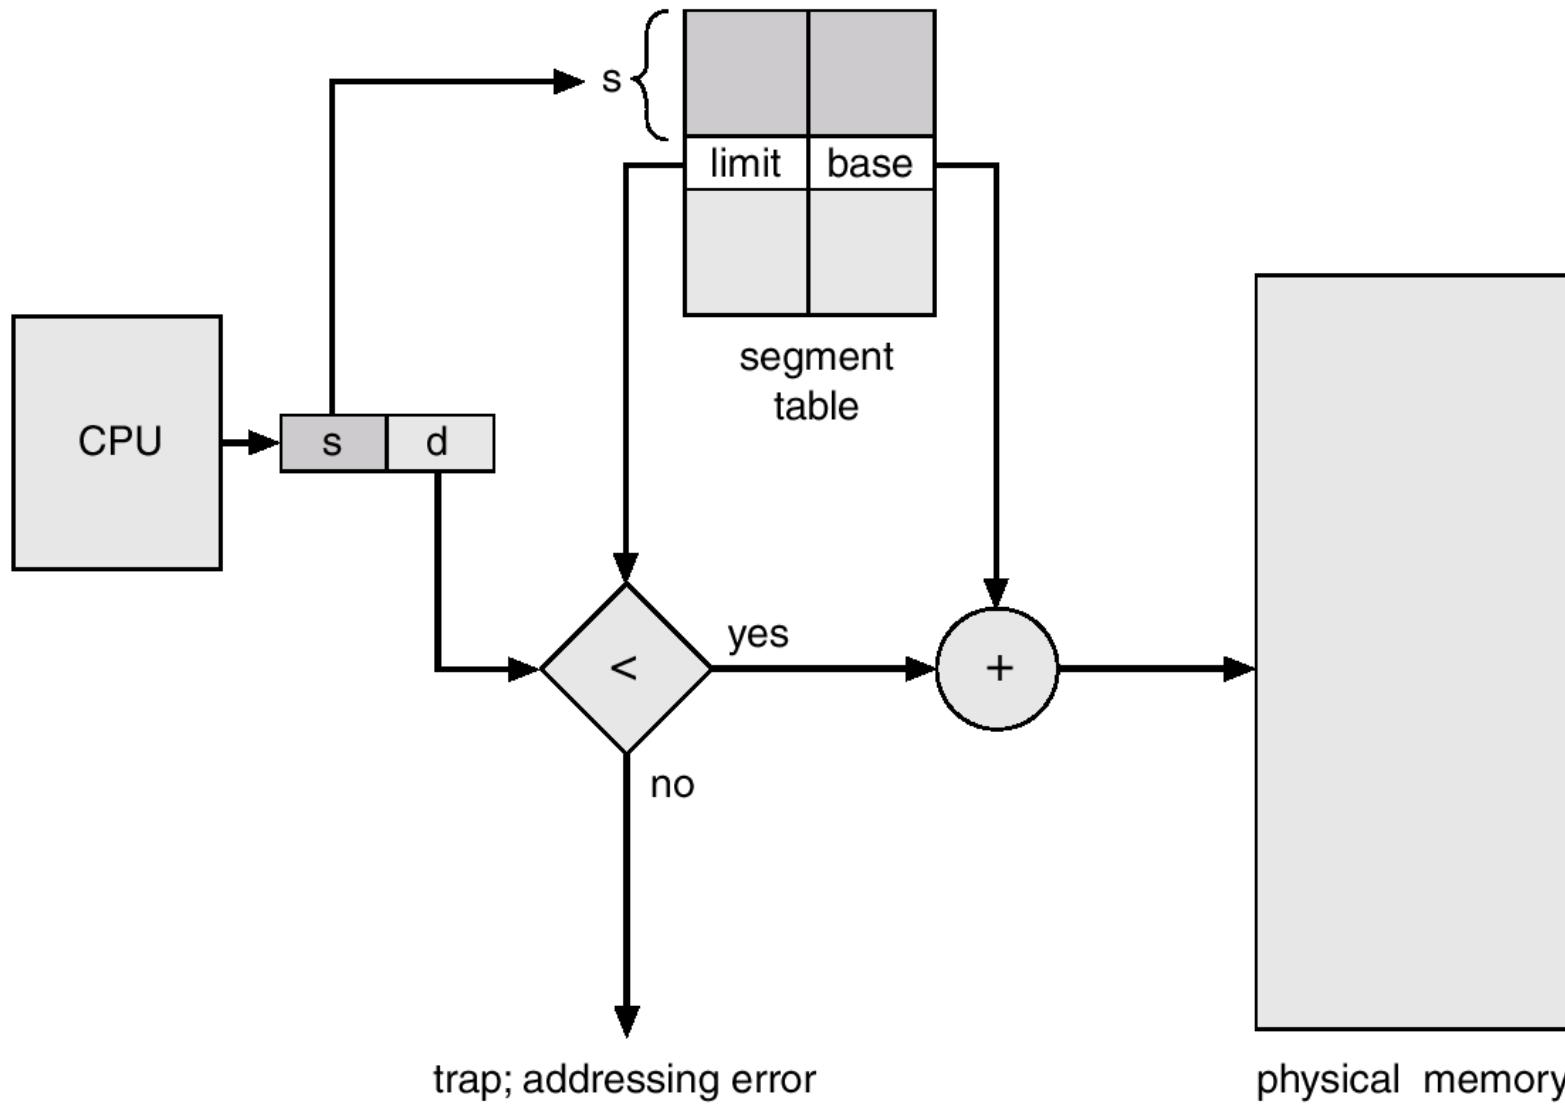
- Logical address consists of a two tuple:
 $\langle \text{segment-number}, \text{offset} \rangle,$

- Segment table – maps two-dimensional physical addresses; each table entry has:
 - ◆ base – contains the starting physical address where the segments reside in memory.
 - ◆ *limit* – specifies the length of the segment.





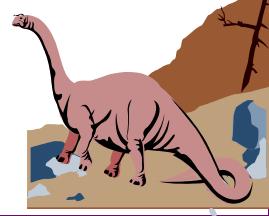
Segmentation Hardware





Segmentation Architecture (Cont.)

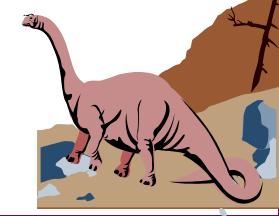
- *Segment-table base register (STBR)* points to the segment table's location in memory.
- *Segment-table length register (STLR)* indicates the number of segments used by a program; segment number s is legal if $s < \text{STLR}$.





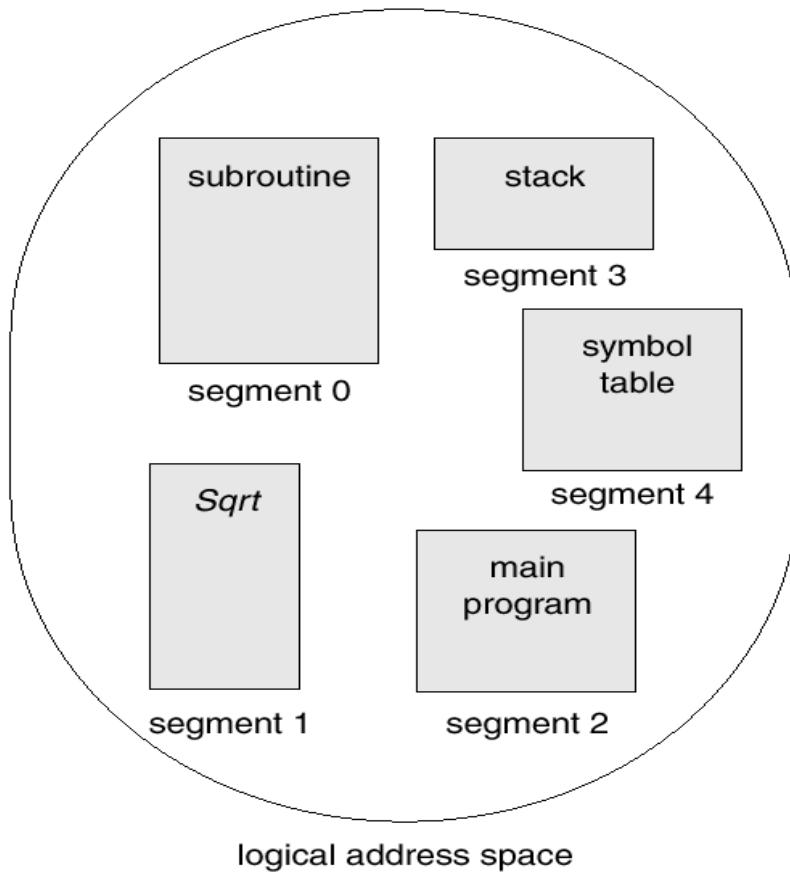
Segmentation Architecture (Cont.)

- Protection. With each entry in segment table, associate:
 - ◆ validation bit = 0 \Rightarrow illegal segment
 - ◆ read/write/execute privileges
- Protection bits associated with segments; code sharing occurs at segment level.
- Since segments vary in length, memory allocation is a dynamic storage-allocation problem.
- A segmentation example is shown in the following diagram



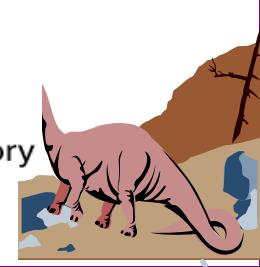



Example of Segmentation



	limit	base
0	1000	1400
1	400	6300
2	400	4300
3	1100	3200
4	1000	4700

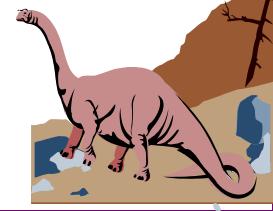
segment table





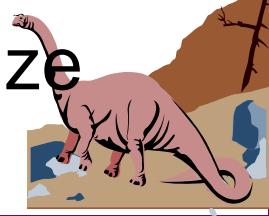
Chapter 9: Memory Management

- Background
- Swapping
- Contiguous Allocation
- Segmentation
- **Paging**
- Advanced Page Table Structure
- Segmentation with Paging





Paging

- Contiguous memory allocation method suffers from the external fragmentation problem
 - Paging method allows logical address space of a process to be noncontiguous; a process is allocated physical memory whenever the latter is available
 - How?
 - ◆ Divide physical memory into fixed-sized blocks called **frames** (帧) (size is power of 2, between 512 bytes and 8192 bytes).
 - ◆ Divide logical memory into blocks of same size called **pages** (页).
- 



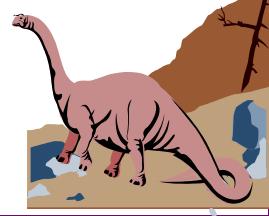
Paging (Cont.)

- Keep track of all free frames.
- To run a program of size n pages, need to find n free frames and load program.
- Set up a page table to translate logical to physical addresses.
- Internal fragmentation.





Address Translation Scheme

- Address generated by CPU is divided into:
 - ◆ *Page number (p)* – used as an index into a *page table* which contains base address of each page in physical memory.
 - ◆ *Page offset (d)* – combined with base address to define the physical memory address that is sent to the memory unit.







logical address $\langle p, d \rangle$



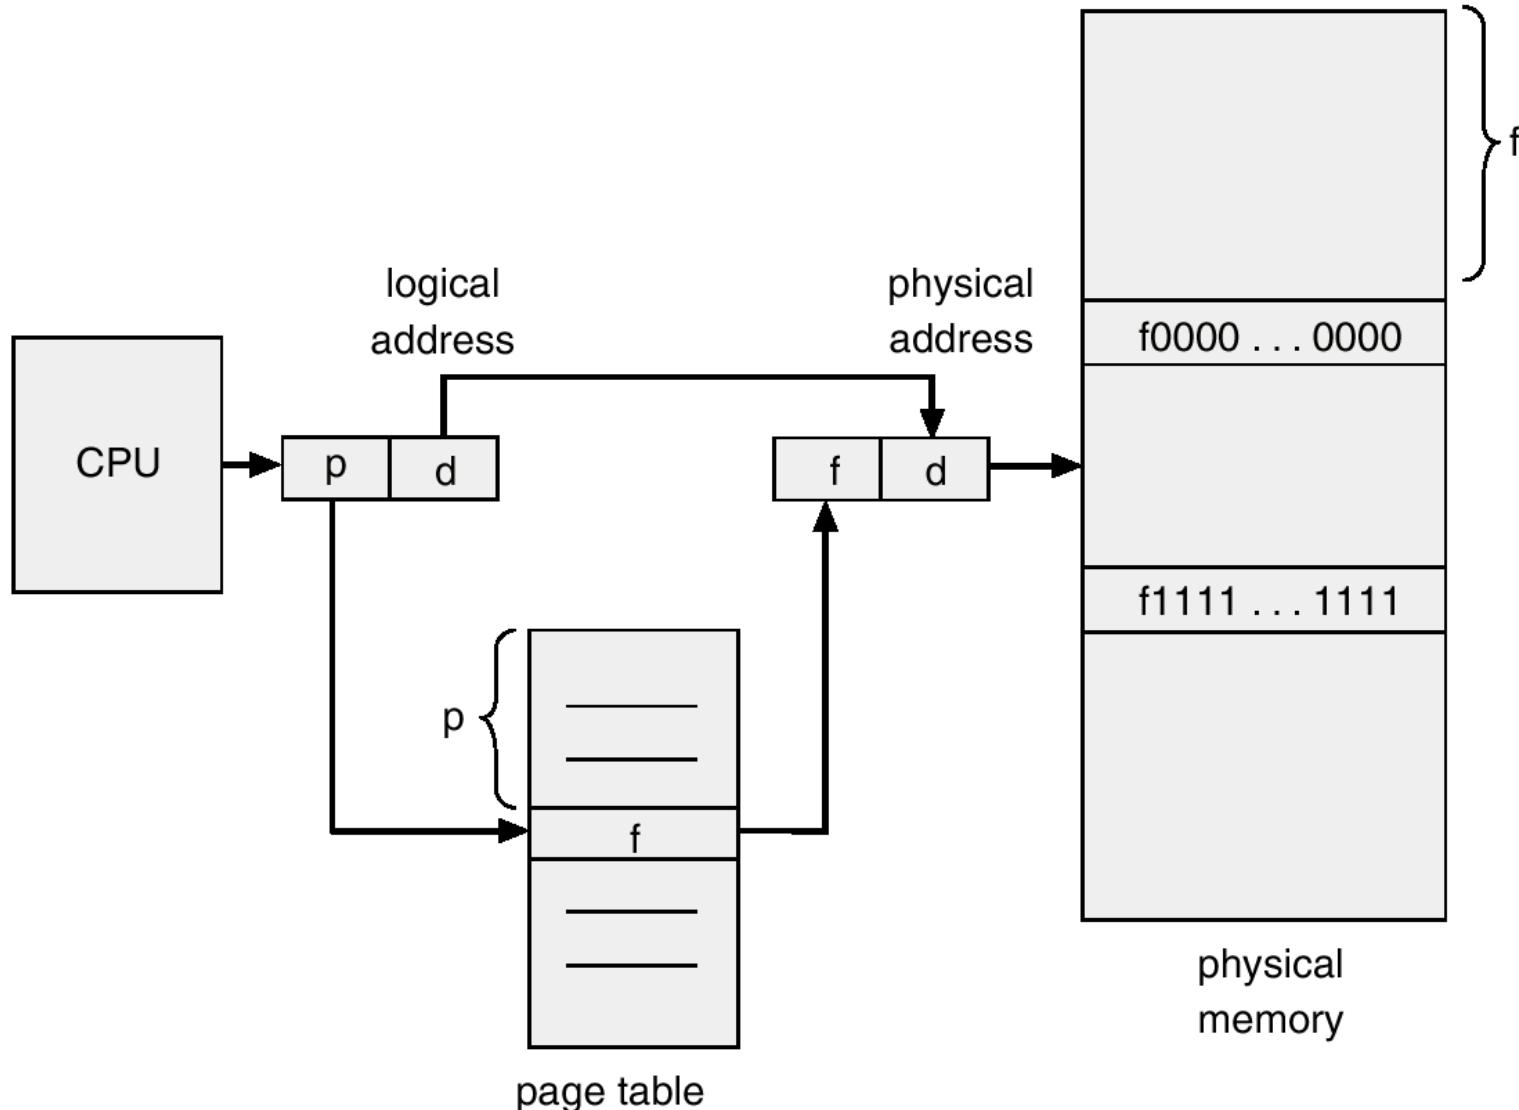
page table

logical address $\langle 1, d \rangle$ translates to physical address $\langle 2, d \rangle$






Address Translation Architecture







Paging Example



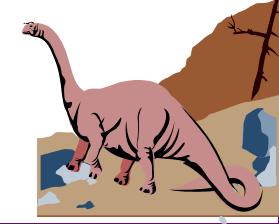
logical
memory



frame
number



physical
memory





Paging Example

0	a
1	b
2	c
3	d
4	e
5	f
6	g
7	h
8	i
9	j
10	k
11	l
12	m
13	n
14	o
15	p


logical memory

0	5
1	6
2	1
3	2

page table

0	
4	i
	j
	k
	l
8	m
	n
	o
	p
12	
16	
20	a
	b
	c
	d
24	e
	f
	g
28	h

physical memory
Southeast University

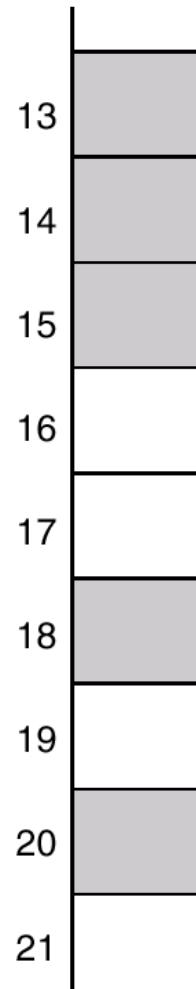
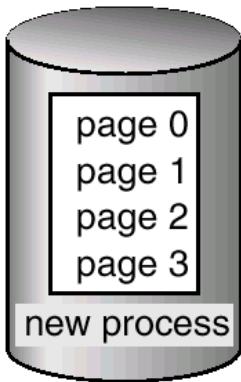




Free Frames

free-frame list

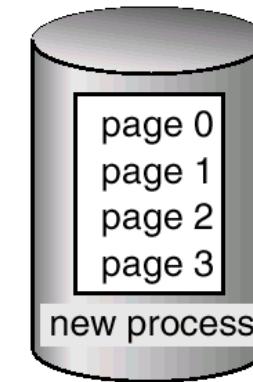
14
13
18
20
15



(a) Before allocation


free-frame list

15



0	14
1	13
2	18
3	20

new-process page table



(b) After allocation



Implementation of Page Table

- Page table must be kept in main memory.
- Question: Why is a page table hard to entirely fit into L2 cache? What will be the size of a page table, if assuming 32 bits virtual address, 4GB physical memory and 4KB page/frame size?

32-bit Virtual Address

20 bits	12 bits
---------	---------

Page # In-page offset

32-bit Physical Address

20 bits	12 bits
---------	---------

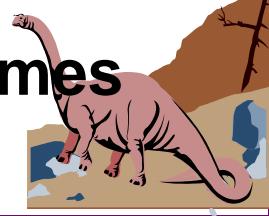
Frame # In-frame offset

◆ 32 bits required to locate a byte in physical mem

- ✓ 4 GB of Physical Memory = 2^{32} bytes.



◆ 20 bits required for frame number.

- ✓ 2^{32} bytes of memory/ 2^{12} bytes per frame = 2^{20} frames



Implementation of Page Table

- Page table must be kept in main memory.
- Question: Why is a page table hard to entirely fit into L2 cache? What will be the size of a page table, if assuming 32 bits virtual address, 4GB physical memory and 4KB page/frame size?




◆ So each page table entry is approximately 4 bytes.
(20 bits frame number is roughly 3 bytes and access control contributes 1 byte)




Implementation of Page Table

- Page table must be kept in main memory.
- Question: What is the size of a page table, if assuming 32 bits virtual address, 4GB physical memory and 4KB page/frame size? Why is a page table hard to entirely fit into CPU cache?

32-bit Virtual Address

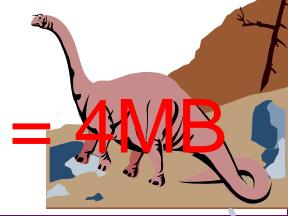


32-bit Physical Address



◆ Page table = 2^{32} bytes / 4KB per page = 2^{20} entries


◆ Page table size = 2^{20} entries * 4B/entry = 4MB





A Quiz

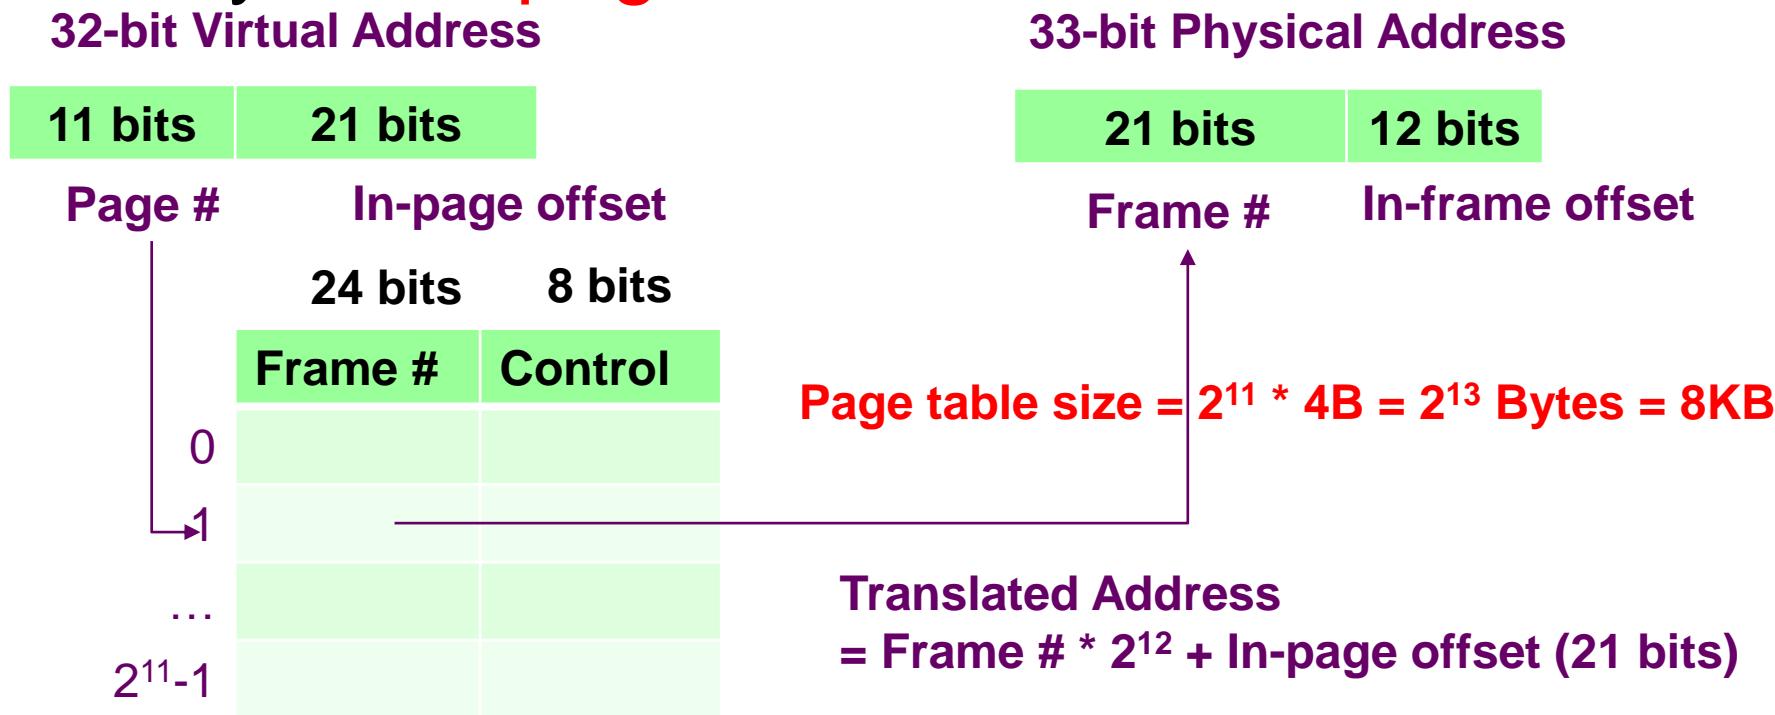
- Q1: What will be the size of a page table, if assuming 32 bits virtual address, 8GB physical memory, **8KB page size**, and 4KB frame size?





A Quiz

- Q2: What will be the size of a page table, if assuming 32 bits virtual address, 8GB physical memory, **2MB page size**, and 4KB frame size?




- Q3: What are the pros & cons of larger page size




Implementation of Page Table

- *Page-table base register (PTBR、页表基址寄存器)* points to page table existing in main memory
- In this scheme every data/instruction access requires two memory accesses:
One for the page table and one for the data/instruction.




- TLB 页表缓存:**
- The two-memory-access problem can be solved by the use of a special fast-lookup hardware cache called *associative memory* or *translation look-aside buffers (TLBs)*



Paging Hardware With TLB (页表缓存)

Translation of virtual address (p, d)

- ◆ If an entry with the key p can be found in the TLB or associative memory, returns the value of frame #
- ◆ Otherwise, get the frame # value from the page table that exists in main memory





TLB based on Associative Memory


■ Associative memory 关联存储器—parallel search

MAPPING from KEY —→ VALUE

Parallel Search {
Entry 1
Entry 2
Entry 3
Entry 4

Page #	Frame #

■ An Example






Effective Access Time

- Associative Lookup = ε time unit
- Assume memory cycle time is 1 time unit
- Hit ratio – percentage of times that a page number is found in the associative registers
- Hit ratio is related to the number of associative registers.

- Hit ratio = α
- Effective Access Time (EAT)



$$\begin{aligned} \text{EAT} &= (1 + \varepsilon) \alpha + (2 + \varepsilon)(1 - \alpha) \\ &= 2 + \varepsilon - \alpha \end{aligned}$$





Memory Protection


- Memory protection implemented by associating protection bit with each frame.
- *Valid-invalid* bit attached to each entry in the page table:
 - ◆ “valid” indicates that the associated page is in the process’ logical address space and is thus a legal page.
 - ◆ “invalid” indicates that the page is not in the process’ logical address space.





Memory Protection (Cont.)

- We can use a *page table length register (PTLR)* that stores the length of a process's page table. In this way, a process cannot access the memory beyond its region.
Compare this with the base/limit register pair.
- We can also add read-only, read-write, or execute bits in page table to enforce **r-w-e** permission.









Advantage of Paging Method: Shared Pages


■ Shared code

- ◆ One copy of read-only (reentrant) code shared among processes (i.e., text editors, compilers, window systems).

■ Private code and data


- ◆ Each process keeps a separate copy of the code and data.
- ◆ The pages for the private code and data can appear anywhere in the logical address space.





Chapter 9: Memory Management


- Background
- Swapping
- Contiguous Allocation
- Segmentation
- Paging
- Advanced Page Table Structure
- Segmentation with Paging





Advanced Page Table Structure

- Hierarchical Paging
- Hashed Page Tables
- Inverted Page Tables






Hierarchical Page Tables

- Why the multiple-level page table is needed?
- **Answer:** A single-level page table may become too big to fit into the physical memory of a commodity machine.


- ◆ Assume a 64-bit computer (which means **64-bit virtual address space**), which has **4KB frames/pages** and **4GB** of physical memory

- ◆ In the single-level page table, 2^{64} addressable bytes / 2^{12} bytes per page = 2^{52} page entries






Hierarchical Page Tables

- ◆ One page table entry contains: Access control bits (like Page present, RW) + Physical page #
 - ◆ **20 bits required for physical page number.**
 - ✓ 4 GB of Physical Memory = **2^{32} bytes**.
 - ✓ 2^{32} bytes of memory/ 2^{12} bytes per page = **2^{20} physical pages**
 - ◆ So each page table entry is approximately **4 bytes**. (20 bits physical page number is roughly 3 bytes and access control contributes 1 byte)
 - ◆ Now page table size = $2^{52} * 4$ bytes = **2^{54} bytes**
 - Hence, the size of single-level page table is **2^{54} bytes (16 petabytes) per process**, which is a very huge amount of memory.
- 




Hierarchical Page Tables

- **A Solution:** Break up the logical address space into multiple page tables.
 - If we **page the page table** too, we can magically bring down the memory required
 - ◆ The first-level page table contains **2^{52} page entries**
 - ◆ If we page the first-level page table, then one page contains $4\text{KB} / 4 \text{ bytes per entry} = 1024 = 2^{10}$ **entries**
 - ◆ So the first-level page table is divided into **2^{42} pages**
 - ◆ So the second-level page table needs **2^{42} entries**
 - ◆
 - ◆ The fifth-level page table only needs **2^{12} page entries**, as low as four pages, just 16 KB memory
- 



Two-Level Page-Table Scheme






Two-Level Paging Example for 32-bit Operating Systems

- A logical address (on a 32-bit machine with 4K page size) is divided into:
 - ◆ a page number consisting of 20 bits, and
 - ◆ a page offset consisting of 12 bits.
- Since the page table itself is also paged, the page number is further divided into:
 - ◆ a 10-bit page number, and
 - ◆ a 10-bit page offset.
- Thus, a logical address is as follows:

page number		page offset
p_1	p_2	d
10	10	12



where p_1 is an index into the outer page table, and p_2 is the displacement within the page of outer page table.





Address-Translation Scheme

- Address-translation scheme for a two-level 32-bit paging architecture





Quiz for Two-level Page Table

■ 某计算机采用二级页表的分页存储管理方式，按字节编制，**页大小为 2^{10} 字节**，页表项大小为2字节。逻辑地址结构为：页目录号、页号、页内偏移量，逻辑地址空间大小为 2^{16} 页，则表示整个逻辑地址空间的页目录表中包含表项的个数是（）

- A、64 B、128 C、256 D、512

页目录号	页号	页内偏移
p_1	p_2	d



■ 答案：B

$$2^{16} / (2^{10} \text{ bytes}/2 \text{ bytes}) = 2^7 = 128$$

■ 页内偏移d占10bits；页表也分页了，每页有 2^9 页表项，所以页号p2占用9bits； $p_1 + p_2 = 16$ bits



Three-level Paging Scheme for 64-bit Operating System




- The 2nd outer page table still needs 16GB memory, since 2^{32} number of table entries \times 4 bytes per table entry = 2^{34} bytes memory







Intel x86-64

- Current generation Intel x86 architecture
- 64 bits address space is ginormous (> 16 exabytes)
- In practice only implement 48 bit addressing
 - ◆ Four levels of paging hierarchy



- ◆ Multiple page sizes of 4 KB, 2 MB, 1 GB





Hashed Page Tables

■ Common in address spaces > 32 bits.

■ Motivation



- ◆ On a 64-bit operating system, the third-level page table is still too large to fit in main memory
- ◆ In a 32-bit or 64-bit address space of a process, most part of it is unused

■ Solution base on Chained Hash Table

- ◆ The virtual page number is hashed into a page table. This page table contains a chain of elements hashing to the same location.
- ◆ Virtual page numbers are compared in this chain searching for a match. If a match is found, the corresponding physical frame is extracted.




Hashed Page Tables







Inverted Page Table

- **Motivation:** All the previous schemes need to maintain a page table for each process.
- One entry for each real page of memory (frame)
- Entry consists of the virtual address of the page stored in that real memory location, with information about the process owning that page
- Decreases memory needed to store each page table, but increases time needed to search the table when a page reference occurs.
- Use hash table to limit the search to one — or at most a few — page-table entries.






Inverted Page Table Architecture





Chapter 9: Memory Management


- Background
- Swapping
- Contiguous Allocation
- Segmentation
- Paging
- Advanced Page Table Structure
- Segmentation with Paging





Example: The Intel Pentium


- Dominant industry chips
- Pentium CPUs are 32-bit and called IA-32 architecture
- Current Intel CPUs are 64-bit and called IA-64 architecture
- Many variations in the chips, cover the main ideas here





The Intel IA-32 Architecture

■ Supports segmentation with paging



■ CPU generates logical address

◆ Selector given to segmentation unit

- ✓ Which produces linear addresses
- ✓ Up to 16K segments per process




s: segment number
g: whether in GDT or LDT
p: protection bits

◆ Linear address given to paging unit


- ✓ Which generates physical address in main memory
- ✓ Paging units form equivalent of MMU



Intel IA-32 Segmentation with Paging




Support two-level page table





Intel IA-32 Architecture Supports Multiple Page Size


Pages sizes can be either 4 KB or 4 MB





Intel x86-64

- Current generation Intel x86 architecture
- 64 bits is ginormous (> 16 exabytes)
- In practice only implement 48 bit addressing
 - ◆ Page sizes of 4 KB, 2 MB, 1 GB
 - ◆ Four levels of paging hierarchy




- Can also use PAE (page address extension) so virtual addresses are 48 bits and physical addresses are 52 bits






Example: ARM Architecture

- Dominant mobile platform chip (Apple iOS and Google Android devices for example)
- Modern, energy efficient, 32-bit CPU
- 4 KB and 16 KB pages
- 1 MB and 16 MB pages (termed **sections**)
- One-level paging for sections, two-level for smaller pages



Example: ARM Architecture (Cont.)




■ Two levels of TLBs

- ◆ Outer level has two micro TLBs (one data, one instruction)
- ◆ Inner is single main TLB
- ◆ Firstly, inner is checked, on miss outers are checked, and on miss page table walk performed by CPU



Concluding Marks

- OS creates, for each process, an illusion of continuous memory address space, based on the paging/segmentation mechanism



- **Question:** Why the mainstream CPU chips organize the page#-to-frame# mapping table in a hierarchical way, instead of using hashed page table or inverted page table?



End of Lecture

附加内容

